
Federated Solar-Flare Prediction on SWAN-SF: A Provenance, Leakage, and Evaluation-Protocol Audit

Deepanshu Balhara¹ and Daves Vats²

Independent Researchers^{1,2}

October 2026

*Journal submission edition — regenerated from the repository record (v4.11); source: `submission/src`,
generator: `tools/build_submission.py`*

Abstract — We audit what happens when federated learning is evaluated on the public SWAN-SF flare benchmark. Six *simulated* clients hold disjoint Dirichlet-skewed shards of the cleaned export under a frozen protocol: immutable splits, validation-only selection and calibration, a runtime leakage audit, and single-pass test evaluation. A provenance audit against the raw benchmark (normalisation-invariant instance matching, 98.7–100% coverage, 100% test-side label agreement) shows the cleaned export’s same-partition train/test pairing shares instances: every flaring test window (6,234) also appears in training, and 85.7–90% of training positives are TimeGAN-synthetic. On a leakage-free fold (training partitions 1–4, single-pass test on partition 5) discrimination generalises for all arms (ROC-AUC 0.906–0.978) and the federated-versus-centralised gap disappears. The deployability verdict, however, is set by the corrected persistence baselines: on TSS, XGBoost (0.848) and logistic regression (0.489) alone clear 24-hour-lagged persistence (0.418), and on HSS no arm clears it (0.434); no arm beats same-region label inertia (TSS 0.967) at this 24-hour label geometry. A BatchNorm-transport diagnostic then shows the in-partition FedAvg collapse (test ROC-AUC 0.575) to be a composite evaluation artefact — untransported normalisation statistics, a mechanism documented in the federated-learning literature since FedBN, compounded by a saturated validation monitor and the checkpoint rule that monitor feeds — rather than a federated-learning failure: the same weights score 0.951 with local statistics, and a controlled no-BatchNorm run leaves both FedAvg and FedProx stable. On the raw, unbalanced substrate the federated MLP collapse is BatchNorm-under-federation-conditional, and two controls establish it: the checkpoint-level BN diagnostic finds that pooled-statistics recalibration — validated as an instrument on the centralised model, where it should reconstruct the model’s own statistics and does preserve the ROC-AUC ranking — collapses every federated arm to chance, so a pooled-statistics transport fix would not rescue federation there, while evaluation-time batch statistics leave the deficit standing; and a no-BatchNorm re-training, changing nothing else, restores the federated MLP to the sequence arms’ ROC-AUC ranking (FedAvg 0.963/0.366, FedProx 0.975/0.345, the latter above the centralised MLP on ROC-AUC; FedProx-LSTM, which also carries no BatchNorm, reaches ROC-AUC 0.970 and PR-AUC 0.404). Calibration, not detection, is the deployability boundary: no validation-fit decision layer survives the prevalence shift. Standard verification metrics (TSS, HSS, Brier skill, persistence and climatology baselines, with reliability for the in-partition arms) accompany every stored operating point of every substrate. This is the first federated evaluation on SWAN-SF and the first instance-level provenance audit of the cleaned export; prior federated flare forecasting (Fu et al., 2023 preprint) and prior random-split leakage warnings (the benchmark’s own creators, 2019) both exist. Every number regenerates from committed result artefacts, and the underlying dataset cache is frozen under a separate 20-file SHA-256 manifest verified by a committed script.

Keywords: federated learning; solar flare prediction; SWAN-SF; benchmark audit; data provenance; evaluation protocol; BatchNorm; non-IID data; class imbalance

Source code and all experimental artefacts: <https://github.com/Daveshvats/Federated-Learning-for-Distributed-Space-Weather-Monitoring>

1 Introduction

Solar-flare prediction from photospheric magnetogram time series is a benchmark-driven machine-learning problem. The community’s public reference is SWAN-SF [1], and the machine-learning literature on it has grown around a small number of evaluation conventions: randomised splits of pooled windows, the cleaned and rebalanced export, and a reporting culture built on ROC/PR summary statistics. Those conventions carry risk. The benchmark’s own creators warned that sliding-window temporal coherence “invalidates the randomness assumption, thus impacting all sampling practices” [2]; subsequent work has echoed the warning [3], and the cleaned export’s release paper explicitly prescribes pairing temporally-preceding training partitions with a later test partition [4]. Machine-learning flare-prediction models built on photospheric vector-magnetogram summaries—the SHARP parameter suite [5] and its descendants—have demonstrated statistically skilful classification of flaring versus non-flaring active regions [6, 7]; the predictive signal lives in quantities such as total unsigned current helicity, magnetic free energy, and flux-emergence proxies. What has not existed, per the documented search protocol of Section 2, is a federated treatment of SWAN-SF—with its cleaned and raw substrates, its leakage structure, and its centralised baselines—under an evaluation protocol that audits its own benchmark artefacts.

Why federate a public benchmark whose data is, in reality, freely poolable? Two honest reasons. Methodologically, cross-silo federated learning [8, 9] is the standard instrument for studying non-IID data heterogeneity under a controlled protocol: it prescribes exactly what moves between data holders (model parameters) and exactly what does not (raw observations), which makes it a clean lens for asking how much of a benchmark’s apparent difficulty survives when the data is deliberately fragmented. Operationally, the SHARP-style parameters at the centre of this pipeline are dominated by a single public instrument (SDO/HMI via JSOC); the six-regional-agency narrative sometimes attached to federated space-weather proposals does not describe this data and is not claimed here. The six clients of this study are a *simulation protocol*—Dirichlet label skew across disjoint shards—not a claim about any institution, and the present study does not claim that any agency’s flare data cannot be pooled.

In cross-silo FL, each data-holder trains a shared architecture locally and transmits model parameters to an aggregation server, which averages them into a new global model; raw measurements do not move. FL has proven transformative in medicine and digital health, where privacy constraints really are binding [10]. Federated flare forecasting itself is not new: Fu et al. posted a federated transfer-learning preprint in July 2023 (Section 2) [11]. As a body of practice, FL also carries its own evaluation-protocol hazards—statistics that travel with weights (most visibly BatchNorm running moments, a failure mode documented since FedBN [12]), drift monitors that saturate, and checkpoint rules that quietly restore the wrong model—and a benchmark audit is exactly the place to expose how those hazards interact with a published comparison.

This paper provides that treatment, and its character is an audit: we ask what a federated pipeline actually achieves on SWAN-SF, what the benchmark’s own artefacts do to the answer, and which parts of the observed behaviour belong to the algorithms rather than to the evaluation machinery. We present a complete, reproducible federated framework—implemented, versioned, and released as open source—that *simulates* six clients (labelled A–F; no real institution participates), each holding a non-IID shard of the benchmark in both its cleaned/rebalanced form and its raw, naturally imbalanced form, and evaluates whether federation can approach the centralised upper bound without any data movement. The framework emits M/X flare probabilities whose calibration, thresholds, and false-alarm rates are themselves audited quantities in this paper—the evaluation finds they do not survive the prevalence shift, and no operational alerting claim is made. The downstream flare→CME→storm→GIC chain is physically far from the quantity modelled here and is not evaluated (Section 8; see [13] for that literature). The specific contributions are:

1. **A problem formulation and system architecture** that maps flare-probability estimation onto a cross-silo federated classification problem with a safety-critical, recall-weighted objective, and that operationalises client heterogeneity as a concrete, simulable protocol (Section 3 and Section 4).
2. **A non-IID federation protocol for SWAN-SF** that combines Dirichlet label-skew partitioning across six simulated regional clients, a server-aware Fed-Focal loss [14] for extreme class imbalance, and optional distribution-aware aggregation and per-client SMOTE over-sampling (both studied as ablation arms; SMOTE is inert on the cleaned variant and an executed negative at natural prevalence), on top of FedAvg [8], FedProx [15], and SCAFFOLD [16] training loops (Section 4).
3. **A frozen, leakage-audited evaluation protocol and a data-provenance audit of the benchmark export.** The protocol comprises an immutable train/validation/test contract, calibration and threshold selection on validation only, single-pass test evaluation, a centralised-MLP architecture control, five-seed statistics, a 24-cell component-ablation grid, region-disjoint split generation, natural-prevalence operating points, untouched client hold-outs, a training-budget audit, event-level alerting metrics, and a calibration-comparison harness (Section 5). The provenance audit aligns every cleaned window to the raw SWAN-SF instances (Harvard Dataverse; normalization-invariant matching) and finds that the cleaned export’s same-partition train/test pairing shares instances—including 100% of the flaring test windows and 85.7–90% TimeGAN-synthetic training positives—so the study re-evaluates all headline claims on a leakage-free temporally-preceding fold (Section 5.7). The leakage phenomenon itself is documented prior art—the benchmark’s creators warned about random-split overlap in 2019 [2] and the cleaned export’s release prescribes temporally-preceding folds [4]; what is new here is the instance-level quantification, performed inside a federated study.
4. **A leakage-free federated evaluation.** On the temporally-preceding fold (training partitions 1–4, single-pass test on partition 5, natural 1.31% prevalence), all arms generalise (ROC-AUC 0.906–0.978); the in-partition federated-versus-centralised gap disappears (FedProx 0.976/0.308 vs architecture-matched centralised MLP 0.973/0.382); the threshold-transfer and calibration failures persist; and an event-level evaluation on 65 physical M/X-class events shows detection is not the deployability boundary—XGBoost’s false-alarm burden is roughly 7× lower than every neural arm’s (Section 6.1).
5. **A raw-substrate re-execution and the BatchNorm-under-federation-conditional collapse.** Retraining the identical frozen protocol on the raw, unbalanced benchmark (natural prevalence, no synthetic positives, the release’s imputation and normalisation reproduced in-pipeline with train-only parameters) separates the arms decisively: every federated *MLP* arm degrades sharply (FedAvg detects zero of 65 events — 38 (58.5%) under the diagnostic’s batch-statistics treatment, the only evaluation with a usable validation operating point, Section 6.3; SCAFFOLD 0.768/0.057) while the same protocol with the sequential SolarLSTM restores the federated standing (FedProx-LSTM 0.970/0.404; SCAFFOLD-LSTM the substrate’s highest event-level detection, 55/65 at 0.4 false-alarm windows per day — an implementation-case result, since the arm runs as the sign-inverted, cold-start variant disclosed in Section 4) — and an architecture-only control re-attributes the split: replacing the MLP’s three BatchNorm layers with the identity and changing nothing else restores FedAvg to 0.963/0.366 and FedProx to 0.975/0.345, at or above the sequence arms’ ROC-AUC ranking (FedProx above the centralised MLP on ROC-AUC), so the collapse is dominated by BatchNorm under federation rather than by the flattened-feature representation. A second seed holds all four sequence arms within ± 0.005 ROC-AUC, and per-client SMOTE at natural prevalence is a clean negative (Sections 6.2–6.4).

6. **A BatchNorm-transport diagnostic and a controlled no-BatchNorm experiment, which re-attribute the in-partition FedAvg-versus-FedProx gap.** The underlying mechanism—weight averaging without normalisation statistics—is established knowledge in the FL literature since FedBN [12] and its successors [17–19]; the contribution here is the composite failure-mode dissection on this pipeline and the self-correction of our own earlier artefact. The shipped implementation transports model parameters only, so BatchNorm running statistics remain at initialisation on the server for the entire run; under that evaluation the FedAvg global model drifts from 0.950 to 0.098 test ROC-AUC over 50 rounds, and a saturated validation monitor ($F_1 = 2p/(1+p)$ whenever every prediction is positive) restores the collapsed round-20 checkpoint as the “best” model. The same weights score 0.951 under correct normalisation and 0.953 even at round 50; with BatchNorm layers removed, both FedAvg and FedProx train stably to near-identical test quality, so the published algorithm gap is substantially an evaluation-implementation artefact (Sections 6.11–6.12).
7. **An in-partition data-locality study, disclosed as such.** On the same-partition protocol (331,185 test windows, 1.88% prevalence) whose instance overlap the audit quantified, FedProx attains ROC-AUC 0.954 (97.7% of the XGBoost reference’s 0.977) against the architecture-matched centralised MLP’s 0.888, a comparison whose optimisation budgets are documented and budget-matched; these numbers are retained as protocol-stability evidence with the corrected attribution above, not as generalisation estimates (Section 6).
8. **A cross-model interpretability audit.** SHAP analysis of both the centralised XGBoost reference and the federated global model, with physical-parameter-group aggregation, shows both concentrate decisions on current helicity, the Schrijver flux-emergence proxy [20], and Lorentz-force components—consistent with established flare physics (Section 6).

We are explicit about what is *not* claimed: per-client SMOTE contributes nothing on either substrate (zero synthetic samples on the cleaned variant; an executed negative at natural prevalence); the transport protocol moves full model weights, not data, and confers no cryptographic privacy (the secure-aggregation sketch verifies arithmetic cancellation only); the six clients are simulated, not real agencies; and the downstream flare→CME→storm→GIC chain is discussed as a boundary, not evaluated.

The released framework also carries a documented revision history (Appendix A): a multi-version stability iteration in which Dirichlet concentration, focal-loss alpha clamping, optimiser choice, and dataset auditing were each adjusted in response to diagnosed failure modes, and the audit findings reported here. We keep that record because reproducible failure diagnosis is transferable know-how for first FL deployments in the physical sciences.

The remainder of the paper is organised as follows. Section 2 reviews flare prediction and federated learning. Section 3 formalises the problem. Section 4 details the system architecture and methodology. Section 5 describes the experimental setup, Section 6 presents and analyses results, and Section 7 discusses operational integration and engineering lessons. Section 8 states limitations and future work, and Section 9 concludes.

2 Background and Related Work

2.1 Solar flare prediction from photospheric magnetograms

The physical premise of data-driven flare forecasting is that the photospheric magnetic field of an active region encodes the stress budget that a future flare will release. Schrijver [20] demonstrated that a single scalar—the R-value, the unsigned flux within high-gradient polarity-separation corridors—separates major-flaring regions from non-flaring ones with striking reliability, and this proxy remains among the strongest single flare predictors identified to date.

Bobra and Couvidat [5] operationalised the machine-learning view of the problem, showing that a support-vector machine trained on 25 SHARP parameters derived from SDO/HMI vector magnetograms can forecast M-class flares with useful skill, and that feature selection matters as much as model choice. Subsequent community-wide blind tests, most notably the multi-method comparison of Leka et al. [6], established two durable findings: (i) no single method dominates across all event definitions and skill metrics, and (ii) the effective ceiling on skill is set by the quality and homogeneity of the training data rather than by classifier complexity.

The European FLARECAST project [7] pushed the operational agenda further, combining multiple forecasting methods and data sources in an open framework, and it explicitly confronted the data-heterogeneity problem that motivates the present paper. On the data side, the SWAN-SF benchmark of Angryk et al. [1] released a multivariate time-series dataset of 4,098 active-region (MVTs) tracks with 24 physical parameters per timestep and flare-history labels, specifically designed to make flare prediction reproducible and comparable; it has since become the community’s reference benchmark, and a family of cleaned and rebalanced variants (random undersampling, Tomek-link cleaning, and TimeGAN-based synthetic augmentation [21]) has been curated to make the benchmark tractable for supervised learning. Recent deep-learning work on SWAN-SF continues to raise the centralised ceiling: contrastive pretraining and multifaceted preprocessing [4] and LSTM-style temporal architectures [22] exploit the sequential structure of the 60-minute observation windows. All of these works, without exception, assume that the full training corpus is available in one place.

2.2 Federated learning under statistical heterogeneity

Federated learning was introduced by McMahan et al. [8] in the FedAvg formulation: clients perform several epochs of local SGD on private data, and a server aggregates the resulting weights by sample-size-weighted averaging. The defining difficulty, identified almost immediately, is statistical heterogeneity: when client data distributions differ—the non-IID setting—local optima diverge from the global optimum, and naive averaging can oscillate or diverge [23]. FedProx [15] addresses this by adding a proximal term $(\mu/2) \|\mathbf{w} - \mathbf{w}^{(t)}\|^2$ to each client’s local objective that anchors local training to the current global model, tolerating partial work and heterogeneous local computation. SCAFFOLD [16] attacks the same problem with control variates that correct client drift using variance reduction. The comprehensive survey of Kairouz et al. [9] situates these algorithms within the broader open-problem landscape, including robust aggregation, personalisation, and privacy. Cross-silo FL—few, reliable, institutionally owned clients—is the deployment regime relevant to space-weather agencies, and it contrasts with the cross-device regime in its stability requirements and long-lived participants.

A distinct and, for this study, decisive strand concerns normalisation statistics under heterogeneity. FedBN [12] established that BatchNorm running statistics are a per-client feature that weight-averaged federation must *not* naively transport, and that keeping them local outperforms both FedAvg and FedProx under feature shift. Wang et al. [17] supplied the convergence analysis of exactly this mismatch, Guerraoui et al. [18] showed corrected statistics restore centralised-level evaluation performance, and BN-SCAFFOLD [19] documents BN drift crippling SCAFFOLD specifically. The parameters-only transport bug we diagnose in Section 6.11 is an engineering-level instantiation of this literature’s central mechanism — the mechanism is textbook; the composite failure chain (initialised-buffer evaluation times a saturated monitor times a checkpoint rule) is the pipeline-specific part.

2.3 Class imbalance in federated learning

Flare prediction inherits an extreme class-imbalance regime: in SWAN-SF’s natural test distribution, roughly 1.9% of windows precede M/X-class flares, while the rebalanced training partitions

sit near parity. The focal loss of Lin et al. [24], originally developed for dense object detection, down-weights easy examples with a $(1 - p_t)^\gamma$ modulating factor and concentrates gradient on hard positives; it has become the default imbalance workhorse. Sarkar et al. [14] adapted focal loss specifically to FL in the Fed-Focal formulation, in which the class-weighting parameter is adjusted using information available at the server about the global class distribution, so that clients with different local positivity rates do not implicitly reweight the global objective. In parallel, synthetic oversampling remains a standard pre-processing response: SMOTE [25] interpolates minority-class neighbours, and mixup [26] creates convex combinations of training pairs. Under federation, rebalancing must be applied locally per client, since the pooled distribution is precisely what no participant can inspect.

2.4 Positioning of this work and search protocol

Two research threads—physics-grounded flare prediction and heterogeneity-robust federated learning—have developed largely independently, and one prior work joins them: Fu, Wan, Huang, Han, and Peng posted “Federated Transfer Learning for Soalr Flare Forecasting” [title sic] to Research Square on 10 July 2023, formatted for *Solar Physics* [11]. That preprint applies federated transfer learning to flare forecasting on its own data substrate; it remains unrefereed, and we found no evidence of formal journal publication as of this revision. It does not evaluate SWAN-SF, reports no provenance audit, and, to our knowledge, no follow-up has appeared. We note that its evaluation protocol—a random 7:1.5:1.5 split of pooled windows (AUC 0.936 on the Tianchi flare dataset)—is precisely the random-split convention whose instance-sharing risk this paper quantifies for SWAN-SF; the comparison is offered as context, not as a score comparison across datasets. Because novelty claims are factual statements about the literature, we document the search protocol behind ours rather than relying on “to the best of our knowledge”. The search was last repeated on 29 September 2026 across general web, arXiv, NASA ADS, and publisher indexes (ScienceDirect, IEEE) reached through a web search engine, using four query strings: (i) *federated learning solar flare prediction*; (ii) *federated learning space weather forecasting*; (iii) *federated learning SWAN-SF benchmark magnetogram*; (iv) *cross-silo federated learning geophysics heliophysics distributed observatories*. Records were included if they apply federated or distributed multi-client learning without raw data pooling to flare prediction or space-weather forecasting; records were excluded if they concern centralised machine learning for flares (the case for all entries of query i), federated forecasting of *solar power or irradiance*, federated *terrestrial* weather forecasting, or federated computing on spacecraft without a scientific forecasting task (the categories returned by query ii). Screening the 30 records returned by the four queries yielded *zero* in-scope records; the raw result listings are committed with the repository (`docs/lit_search/`, each listing embedding its query date since v4.6). That screening *missed* the Fu et al. preprint above—exactly the grey-literature coverage limitation this section flags—and the preprint surfaced only in a targeted verification pass during the present revision. We note two further limitations of this protocol: web-engine coverage of preprint servers is incomplete, and the screening was single-reviewer, so a PRISMA-style dual-screened log remains queued repository work. Three facts bound the novelty claim honestly. The leakage phenomenon is documented prior art: the SWAN-SF creators warned that temporal coherence invalidates randomised sampling [2], Ahmadzadeh et al. repeated the warning [3], and the cleaned export’s release paper prescribes temporally-preceding folds [4]; the BatchNorm-transport mechanism is textbook since FedBN [12]. Within those limits, the present work is the first federated evaluation on the SWAN-SF benchmark (in both its cleaned and raw substrates, against centralised baselines of matched architecture), the first *instance-level* provenance audit of the cleaned export (normalisation-invariant matching of every window to the raw benchmark), and the first leakage-free re-evaluation of federated flare prediction under the benchmark’s intended protocol. It also contributes a rarely documented artifact: an honest engineering log of

the failure modes encountered when FL meets a severely imbalanced physical-science benchmark (Appendix A).

3 Problem Formulation

3.1 Federated risk minimisation

Let K denote the number of participating clients (simulated shards of the pooled benchmark), each holding a private dataset $\mathcal{D}_k = \{(x_i, y_i)\}_{i=1}^{n_k}$ of active-region observation windows $x_i \in \mathbb{R}^d$ and binary labels $y_i \in \{0, 1\}$ indicating whether an M- or X-class flare occurs within the forecasting horizon. Let $\mathcal{D} = \bigcup_{k=1}^K \mathcal{D}_k$ with $n = \sum_k n_k$. The goal of standard centralised learning is to solve

$$\min_{\mathbf{w} \in \mathbb{R}^p} F(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\mathbf{w}}(x_i), y_i), \quad (1)$$

where $f_{\mathbf{w}}$ is the prediction model, ℓ a per-sample loss, and p the number of parameters. In the federated setting this objective is decomposed as a weighted sum of local empirical risks

$$\min_{\mathbf{w}} F(\mathbf{w}) = \sum_{k=1}^K \rho_k F_k(\mathbf{w}), \quad F_k(\mathbf{w}) = \frac{1}{n_k} \sum_{(x_i, y_i) \in \mathcal{D}_k} \ell(f_{\mathbf{w}}(x_i), y_i), \quad (2)$$

with aggregation weights $\rho_k \geq 0$, $\sum_k \rho_k = 1$. The vanilla FedAvg choice is $\rho_k = n_k/n$; Section 4 describes the distribution-aware weights used in this work.

3.2 The data-locality constraint

The operational constraint that distinguishes cross-silo FL from distributed optimisation is not an algorithmic preference but an institutional invariant: for every round t and client k , the server may observe only the weight increment $\Delta \mathbf{w}_k^{(t)} = \mathbf{w}_k^{(t+1)} - \mathbf{w}_k^{(t)}$ (or equivalently the locally trained weights), never x_i or y_i for any sample in $\mathcal{D}_k$. Formally, the communication transcript

$$\mathcal{T} = \{\Delta \mathbf{w}_k^{(t)}\}_{k=1, \dots, K; t=1, \dots, T} \quad (3)$$

is the only quantity that crosses the client boundary. This is the data-locality constraint that motivates cross-silo federated learning in settings where shards genuinely cannot be pooled—proprietary archives, cross-border data governance, organisational silos. We stress that no such constraint binds this benchmark: the SWAN-SF data is public and freely poolable, and the federated protocol here is a controlled simulation of the constraint, not a response to it. The released framework nonetheless enforces the property by construction: the client-side pipeline never serialises raw shards. We further stress the scope of the property: it is *data locality*, not cryptographic privacy. The transcript $\mathcal{T}$ itself can leak information about training data (Section 4.9); defending against an adversary who observes it requires secure aggregation or differential privacy, which are implemented but not evaluated for accuracy cost here.

3.3 Statistical heterogeneity

In genuine cross-silo deployments, participating organisations sample different populations under different label distributions, so the local class priors $p_k = \Pr(y = 1 \mid k)$ differ from the pooled prior $\bar{p}$ —statistical heterogeneity, the defining difficulty of federated optimisation. Because SWAN-SF is a single-instrument, whole-disk benchmark, no physical mechanism produces such skew here; we therefore *impose* it by partitioning the pooled training data among K clients with

a symmetric Dirichlet distribution of concentration α (Section 4), so that the realised local priors span a wide range around $\bar{p}$. The partition is a simulation instrument — a test of federation under heterogeneity, not a claim about how the data is produced: the labels and windows come from one observatory pipeline for every shard. The consequence for optimisation is well documented [15, 23]: local objectives F_k become misaligned, and plain averaging of locally optimised weights introduces drift that grows with the number of local epochs. FedProx counteracts this drift by locally solving

$$\min_{\mathbf{w}} F_k(\mathbf{w}) + \frac{\mu}{2} \left\| \mathbf{w} - \mathbf{w}^{(t)} \right\|^2, \quad (4)$$

which keeps client updates in a trust region around the global model. The strength of this anchor, μ , is a central hyperparameter of this study.

3.4 A safety-critical, recall-weighted objective

Rare-event forecasting inverts the usual classification cost structure: at 1–2% positive prevalence an all-clear forecast is 98–99% accurate and useless, while a missed positive is the event the forecast exists to catch. Following rare-event evaluation practice, we therefore evaluate with the F_β measure

$$F_\beta = \frac{(1 + \beta^2) P R}{\beta^2 P + R}, \quad \beta = 2, \quad (5)$$

where P and R are precision and recall; $\beta = 2$ weights recall twice as heavily as precision. Because the deployed decision rule is a threshold τ on the predicted flare probability, the framework treats τ as a free, per-model parameter selected on held-out validation evidence to maximise Eq. (5), rather than fixing $\tau = 0.5$ by convention. Section 6 shows this choice is not cosmetic: at the conventional threshold, every federated model trained on rebalanced local data is operationally unusable, and threshold optimisation is what converts a well-ranked model into a usable detector.

4 System Architecture and Methodology

4.1 Architecture overview

Figure 1 shows the complete system. Six simulated clients—labelled neutrally A–F (plain identifiers; the labels carry no geographic or institutional mapping, and no real institution participates in this study)—each hold a private, non-IID shard of the benchmark. In every communication round, each selected client trains a local copy of the shared model on its shard (with an imbalance-robust loss, described below), and transmits its model parameters to the aggregation server (the implementation transports full weight tensors both directions; delta/weight-increment encoding and quantised uplinks are the communication-cost model’s future compression path, not what the code does). The server forms the next global model by size-weighted averaging, broadcasts it to all clients, and the cycle repeats for $T = 50$ rounds. Downstream of federation, the global model emits flare probabilities; the evaluation of this paper finds that its calibration and frozen operating points do not survive the prevalence shift, so no alerting claim is made. Any downstream use—CME propagation and storm forecasting, geomagnetic-hazard estimation, GIC and power-system response modelling—would require components this paper does not evaluate, and the flare→CME→storm→GIC chain is far from the quantity modelled here (Section 8). The division of labour is deliberate: everything that touches raw magnetogram data (loading, local training) executes inside the client boundary; everything the server touches is a weight tensor.




Figure 1: Federated benchmark-audit architecture. Six simulated clients (neutral labels A–F; no real institution participates) hold private non-IID shards of the benchmark; only model weight tensors w_k cross the client boundary (the figure and the communication-cost model describe full-weight transport; delta encoding is listed future compression work). The server applies size-weighted FedAvg/FedProx aggregation by default (the distribution-aware weighting $\rho_k \propto n_k \sqrt{\phi_k}$, capped at $2/K$, is studied as an ablation arm). The global model emits flare probabilities; the evaluation finds its calibration and frozen operating points do not survive the prevalence shift, so no alerting claim is made, and no downstream pipeline is evaluated.

4.2 Data: the SWAN-SF benchmark and its cleaned variant

All experiments use the SWAN-SF benchmark [1] (doi:10.7910/DVN/EBCFKM), a multivariate time-series dataset of active-region observation windows drawn from SDO/HMI vector magnetograms, with 24 physically motivated parameters per timestep sampled over 60-minute sliding windows and labels indicating the occurrence of $\geq M$ -class flares within the prediction horizon. The 24 parameters summarise magnetic-field complexity and include, among others: total unsigned current helicity (TOTUSJH), total photospheric magnetic free energy (TOTPOT), total unsigned vertical current (TOTUSJZ), total unsigned flux (USFLUX), the sum of absolute net-current-per-polarity (SAVNCPP), absolute net current helicity (ABSNJZH), mean shear angle and the fraction of high-shear area (MEANSHR, SHRGT45), the Schrijver flux-emergence proxy (R-value), and the three components of the total Lorentz force (TOTFX, TOTFY, TOTFZ) [5, 20].

We ingest the benchmark through its community-maintained cleaned variant¹, combining all five data partitions. The training partitions are class-balanced by this preprocessing (pooled positive rate $\bar{p} \approx 0.489$), while the untouched test partitions retain the natural, severely imbal-

¹Cleaned SWAN-SF Dataset, <https://github.com/samresume/Cleaned-SWANSF-Dataset>, which applies per-partition preprocessing: random undersampling with Tomek links and TimeGAN-based synthetic augmentation for the training partitions [21], LSBZM normalisation, and FPC-KNN imputation.

anced prevalence of approximately 1.9% flare windows across roughly 3.3×10^5 held-out samples. This train–test contrast is not an artefact we introduce; it is inherited from the benchmark’s design philosophy, and it has a consequence that shapes the entire evaluation: models trained at $\sim 49\%$ prior must detect events at a $\sim 2\%$ prior, so absolute probability calibration degrades and decision thresholds must be re-optimised post hoc (Section 4.8).

The raw-benchmark experiments of Sections 6.1 and 6.2 additionally ingest the unprocessed SWAN-SF partitions directly. The raw substrate cannot be trained on as-is: 8.87% of its values are missing, the missingness is structured (the Lorentz-force component TOTFZ is unobserved in 98.7% of training windows), and the 24 parameters span eleven orders of magnitude. We therefore reproduce the cleaned release’s two preprocessing stages in-pipeline from their published description (ApJS, 10.3847/1538-4365/ad7c4a) rather than inheriting their output: a fast Pearson-correlation-based k-NN imputation (within-window linear interpolation for partial gaps; inverse-distance-weighted transfer of the $k=5$ nearest training windows’ series, searched in the space of the five most-correlated observable features, for fully-missing series), and the LSBZM normalisation chain (robust quantile clip, positivity shift, a log/sqrt/Box–Cox stabiliser dispatched per feature by fitted λ , z-score, min–max), with every parameter fitted on the training pool only and a per-feature power-of-two pre-scaling for numerical safety. The reproduction is verified against the released export on the audit-matched windows (observed-position rank agreement in Section 6.2); where the two pipelines differ substantively (the release re-imputes the R-value zero mass as if missing; we preserve zeros as physical “no flux emergence” values), we preserve the raw semantics. The construction, the runner, and the verification are committed as `experiments/raw_substrate.py`, `experiments/run_raw_substrate.py`, and `experiments/run_event_level_raw.py`.

4.3 Client-side feature representation

The framework supports two representations of the same 60-timestep observation windows. In the MLP path used for the reported run, each three-dimensional window $X \in \mathbb{R}^{60 \times 24}$ is reduced to a 144-dimensional statistical vector

$$\tilde{x} = [\mu, \sigma, \max, \min, \delta, s] \in \mathbb{R}^{144}, \quad (6)$$

where each block is applied per base feature: μ and σ are the temporal mean and standard deviation, $\max$ and $\min$ the window extrema, $\delta = X_T - X_1$ the net trend, and s the least-squares slope against normalised time. This six-statistic extraction preserves substantially more temporal information than a temporal mean (peaks, volatility, and rate of change are precisely the flare-relevant quantities) while remaining compatible with weight-averaged federation and with the centralised baselines. In the LSTM path, the raw $\mathbb{R}^{60 \times 24}$ tensor is consumed directly by a recurrent model; this path was executed federated on the raw substrate for all four arms plus the pooled comparator (Section 6.5), with its cleaned-fold pairing remaining future work (Section 8).

4.4 Non-IID partitioning into simulated clients

To emulate label heterogeneity, the pooled training set is partitioned by label-aware Dirichlet sampling: for each class c , a proportion vector is drawn from $\text{Dirichlet}(\alpha, \dots, \alpha) \in \mathbb{R}^K$ and each client receives the corresponding fraction of that class. The concentration α controls heterogeneity strength—small α produces near-exclusive class ownership; large α approaches IID. Figure 2 shows the deterministic realisation used by every experiment in this paper ($\alpha = 1.0$, seed 42) on the class-balanced pool: realised local flare rates span 28.6% to 80.5% and shard sizes range from 2.3% to 35.6% of the pool. This is a deliberately hostile federation: some clients are flare-rich and tiny, others flare-poor and large, mirroring the real asymmetry a cross-silo federation faces




Figure 2: Non-IID partitioning of the pooled, class-balanced training data into six simulated clients by label-aware Dirichlet sampling (concentration $\alpha = 1.0$, RNG seed 42, Section 4); proportions are read directly from the frozen run’s cached client assignment. (a) Realised relative shard sizes. (b) Realised local flare rates against the pooled rate of 48.9%: label skew spans 28.6%–80.5%, emulating the label heterogeneity a federation would face if the shards sampled very different active-region populations (no real institution or region participates).

when some participants’ archives dominate the pooled corpus and others sample much smaller populations. The concentration $\alpha = 1.0$ itself is the outcome of a documented tuning process (Appendix A): early experiments at $\alpha \leq 0.5$ produced pathological partitions (0.3%–99.8% local flare rates, re-measured under the frozen seed) under which no aggregation scheme remained stable.

4.5 Client models

The primary federated architecture is a four-layer multilayer perceptron, $\text{SolarMLP} : \mathbb{R}^{144} \rightarrow \mathbb{R}$, with hidden widths $144 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 1$, batch normalisation, ReLU activations, and dropout 0.3 after each hidden layer; the single output is a raw logit consumed by the sigmoid-based losses below. MLPs are natural federation citizens because every parameter is a numerically meaningful averageable tensor. The released framework also implements SolarLSTM, a (optionally bidirectional) two-layer LSTM with hidden size 128, Xavier/orthogonal initialisation, and a forget-gate bias of one for long-horizon memory, followed by a two-layer classification head [22, 27]. Local training uses AdamW [28] with cosine-annealing warm restarts and gradient-norm clipping at 5.0. We deliberately keep XGBoost out of the federation loop: gradient-boosted tree ensembles store discrete structure that cannot be meaningfully averaged, so the centralised XGBoost serves as the pooled upper bound rather than a federated participant.

4.6 Federation algorithms and aggregation

Three cross-silo algorithms are implemented. **FedAvg** [8]: each round, every client initialises from the broadcast global model, performs $E = 10$ local epochs, and the server averages the returned weights. **FedProx** [15]: identical to FedAvg except each client’s local objective is regularised by the proximal term of Eq. (4) with $\mu = 0.01$, penalising drift from the global model—the mechanism we hypothesise is decisive under the label skew of Figure 2. **SCAF-FOLD** [16]: variance reduction via client and server control variates; it is implemented with the SGD-based update required for exactness of the control-variate correction (AdamW’s adaptive steps would invalidate the correction terms). Five implementation details depart from

the reference algorithm and are disclosed for exactness: (i) the control-variate correction is damped, rescaled to 10% of each parameter’s gradient norm (the undamped correction destabilised training on this benchmark’s loss scale); (ii) gradients are clipped to max-norm 1.0; (iii) the shipped optimiser is SGD with momentum 0.9 at a learning rate of 2.5×10^{-4} — half the global rate — a departure of a milder class, since momentum biases the average-gradient estimate that the control-variate update divides by (disclosed at v4.5); (iv) the client control-variate update is *sign-inverted* relative to Eq. 11 of the reference algorithm: the repository computes $c_i^+ = c_i - c + (y_i^+ - y_i)/(\eta K)$ — the negative of the client’s average local gradient — where Karimireddy et al. specify $c_i^+ = c_i - c + (y_i - y_i^+)/(\eta K)$; a scalar one-client simulation of the code path confirms the inversion exactly (the update evaluates to -1.0000 times the average local gradient), and the wrongly-signed variate estimates push the correction term $c - c_i$ the wrong way over rounds — the damping of item (i) and the variate clamping bound the damage, which is why the arm still trains (found by Dossier R-FS9-R6, disclosed at v4.6); and (v) the arm is trained *cold*: on every path that produced a published operating point — the raw-substrate MLP and LSTM arms — the model starts from a fresh random initialisation at half the global learning rate, not from a converged FedAvg model (the FedAvg-initialised warm start claimed at v3.9–v4.5 exists only on the cleaned-fold path, where the arm is never executed; the v4.6 correction follows the call sites and the queue log, whose validation F1 is pinned at zero through round thirty — the cold-start signature). A cold start at half the global rate is a *harder* protocol than the FedAvg-initialised refinement previously claimed, so the arm’s published trajectory should be read as a from-scratch run. One further arm-level asymmetry is disclosed at v4.5: the published SCAFFOLD operating points were trained before the imbalance factor’s global-prevalence fix (founding-audit finding B13) was extended to the SCAFFOLD local path, so that arm’s focal weight used the hardcoded cleaned-substrate prevalence 0.4887 as $\bar{p}$ rather than the computed global rate the FedAvg and FedProx arms use; on the roughly two-percent-positive raw substrate the imbalance factor saturates at its clamp either way, which bounds the effect. The call site is aligned at v4.5, so all three federated arms share one loss definition in the released code, and the same alignment is extended at v4.6 to the centralised comparators (the pooled MLP, the budget-matched re-run, and the pooled LSTM), whose published operating points likewise predate the plumbing and trained under the same hardcoded-prevalence fallback — an asymmetry of exactly the bounded class disclosed above for the SCAFFOLD arm, disclosed here for the central side as well (Dossier R-FS9-R6, R8-5). Because every published SCAFFOLD operating point was produced by this implementation, the arm’s comparisons against its siblings measure *this implementation*, not the reference algorithm: the raw-substrate cells of Sections 6.2 and 6.5 — the MLP instance’s hardest-failure mode and the sequence instance’s highest event-level detection — are implementation-case evidence for a sign-inverted, cold-start, damped SCAFFOLD, and a sign-corrected, warm-started vanilla replication under the frozen protocol is queued owner-GPU future work whose outcome may differ. The arm was held out of the cleaned-substrate headline configuration for stability reasons and is not run on the cleaned fold, but it is executed on the raw substrate for both encoders, where its MLP instance is the hardest failure and its sequence instance completes the four-arm federation matrix (Sections 6.2 and 6.5).

Aggregation is *plain size-weighted averaging* in the frozen headline configuration ($\rho_k = n_k/n$, the vanilla FedAvg choice); distribution-aware weighting (DA-FL), in which the server weights client k ’s update by

$$\rho_k \propto n_k \sqrt{\phi_k}, \quad \phi_k = \frac{p_k}{\bar{p}}, \quad \rho_k \leq \frac{2}{K}, \quad (7)$$

with p_k the client’s local positive rate and $\bar{p}$ the pooled rate known to the server, is implemented as an *ablation arm* (Section 6.15). The square-root damping and the $2/K$ cap prevent a flare-rich client from dominating the global model while still up-weighting the clients most likely to contain decision-relevant minority samples. Class-rate metadata is deliberately much coarser than raw data and is treated as shareable under the data-locality constraint.

4.7 Fed-Focal loss and local rebalancing

Each client trains with a server-aware focal loss in the Fed-Focal family [14],

$$\ell_{\text{fedfocal}} = -\alpha_t (1 - p_t)^\gamma \log(p_t), \quad \gamma = 2.0, \quad (8)$$

where p_t is the model’s probability of the correct class and α_t interpolates between α and $1 - \alpha$ by label. The positive-class weight α is computed per client as the base value 0.25 [24] scaled by a square-root-damped imbalance factor $\min(\sqrt{\bar{p}/p_k}, 1.5)$ — with the client rate floored at $\max(p_k, 0.01)$ so a flare-free client cannot divide by zero — and a mild round-progress factor in $[0.9, 1.1]$ (linear in round progress; the $[0.7, 1.3]$ range of an earlier revision inflated α past its clamp), then clamped to $(0.05, 0.25)$. Both constants are v4.6 disclosures completing the loss definition (Dossier R-FS9-R6 R8-12). This clamping is a hard-won design decision: unclamped adaptive weighting (earlier versions allowed α up to 0.95) systematically inflated positive-class weight on flare-poor clients, producing degenerate always-flare classifiers with recall ≈ 1.0 and precision ≈ 0.03 (Appendix A). NaN-safe clamping of p_t and of the loss itself guards the pipeline against logit explosions in early rounds. The framework additionally supports per-client SMOTE oversampling [25] and mixup [26]; both are *disabled* in the frozen configuration because the cleaned benchmark’s training partitions arrive pre-balanced (48.9% positive), which renders SMOTE inert—the configured oversampling ratio produces zero synthetic samples. The balancing question is therefore inapplicable to this benchmark variant, and we report it as a null result rather than a contribution (Section 6.15); on the raw substrate, where natural prevalence gives it work to do, the per-client SMOTE ablation has since been executed and is likewise a clean negative (Section 6.6).

4.8 Threshold protocol for safety-critical recall

Because federated training operates on rebalanced local data while deployment operates at the natural 1.9% prior, the emitted probabilities are deliberately miscalibrated, and a fixed 0.5 threshold is meaningless. The frozen protocol therefore performs all selection on validation data, in two stages. First, the decision threshold is selected by sweeping τ over $[0.05, 0.95]$ in steps of 0.005 and choosing, per model, the τ^* that maximises F_β with $\beta = 2$ (Eq. (5)) on the validation set; τ^* is then frozen and applied to the test set exactly once. All reported precision, recall, F1, and accuracy figures are taken at each model’s frozen threshold. Second—because balanced-prevalence validation turns out not to support threshold transfer at all (Section 6.13)—the framework additionally freezes *fixed-FPR operating points*: for each target false-positive budget $f \in \{0.5\%, 1\%, 2\%, 5\%\}$, the threshold is set to the $(1 - f)$ -quantile of *validation negative scores*, a quantity that uses no test information whatsoever; the realised test behaviour at each frozen threshold (achieved FPR, recall, precision, alert rate) is reported in the same single test pass. The negative-quantile rule is numerically validated in the released test suite, and the raw-substrate retraining of Section 6.2 supplies the natural-prevalence validation split that is the ideal substrate for this protocol. ROC-AUC and PR-AUC, being threshold-free, are reported independently.

4.9 Threat model and security posture

The verified property of this framework is data locality (raw data never leaves the client), not end-to-end security. To make that scope auditable, Table 1 states the threat model explicitly: which adversaries are considered, which defences the released code provides, and which are documented but unevaluated.

For the benchmark study reported here—six simulated clients, public data, no real adversaries—none of these defences is enabled, so the reported accuracy numbers are those

Table 1: Threat model of the released framework. “Mitigated in code” means implemented and covered by the integrity test suite; “documented” means design guidance exists but no quantitative evaluation has been performed on this benchmark.

Adversary / attack	Defence	Status
Honest-but-curious aggregation server (observes all weight updates)	Secure-aggregation module (Bonawitz-style masks; tests verify arithmetic round-trip cancellation only — the sketch uses deterministic, publicly reproducible mask keys, so it confers no cryptographic privacy as committed)	Implemented as a numerical sketch; not wired into the training loop; accuracy overhead not measured
Gradient inversion / reconstruction from the transcript $\mathcal{T}$ [29]	Secure aggregation (masks updates); differential-privacy hooks	Documented; not evaluated
Membership inference against a specific client’s training data	Differential-privacy hooks (noise on client updates)	Documented; not evaluated
Malicious client: model poisoning / backdoor insertion	None at present; Byzantine-robust aggregation is a design note	Not addressed
Client impersonation / spoofed updates	Transport security (TLS) and client authentication are deployment infrastructure outside the learning loop	Documented (deployment guidance)
Network interception of updates	TLS in transit	Documented (deployment guidance)

of the undefended pipeline. Any cross-agency deployment would need to re-measure accuracy under secure aggregation and differential privacy before the numbers in Section 6 could be quoted as expected operational performance (Section 8).

4.10 Interpretability layer

To audit what the models learn, SHAP values [30] are computed for the centralised XGBoost reference over a 500-sample test subset with the TreeExplainer and for the federated model with DeepSHAP (100-sample background), yielding mean absolute SHAP attributions per feature. The statistical representation of Eq. (6) carries a 144-entry name contract—each feature is named `<statistic>_<parameter>` (e.g. `mean_TOTUSJH`, `max_R_VALUE`), with the statistic blocks ordered μ (0–23), σ (24–47), $\max$ (48–71), $\min$ (72–95), δ (96–119), s (120–143) over the 24 base parameters—and the improved pipeline preserves these names end-to-end, so the SHAP layer reports physically meaningful attributions directly (the v2.x pipeline lost the mapping and required an analytical decoding pass; Appendix A). Section 6 presents the raw ranking, its statistic-level structure, and the physical grouping used for cross-model comparison.

5 Experimental Setup

5.1 Frozen evaluation protocol

All experiments in this revision run under an immutable evaluation contract, enforced in code before any training is permitted:

- **Splits.** The benchmark’s five training partitions (97,764 windows, balanced, 48.87% positive) form the training pool; the five test partitions (331,185 windows, natural prevalence, 1.88% positive) are the held-out test set and are loaded exactly once at the very end of the pipeline. A stratified validation set (16%, $n=15,643$) is carved from the *training* pool with the root seed (42) and is the only set used for any model selection.
- **Selection isolation.** Hyperparameter monitoring, the best-checkpoint rule, calibration fitting, the F_β threshold search, and the fixed-FPR operating-point thresholds all operate on validation. The test set participates in no selection of any kind (the v2.x pipeline selected thresholds directly on test; this revision removes that violation).
- **Leakage audit.** A runtime audit verifies sample-disjointness of train/validation/test, disjointness and coverage of the client shards, and flags the known limitation that active-region identities are not available in the cleaned export, so region-level separation across the train/validation boundary cannot be guaranteed (Section 8). The pipeline refuses to train if the audit fails.
- **Single-pass evaluation.** Each model’s frozen pipeline (calibrator + threshold) is applied to the test set once; all reported test metrics are recomputed from those predictions (the v2.x code reported a stale accuracy for federated models).
- **Protocol extensions released with this revision.** The following components are implemented, unit-tested (the integrity battery of `tests/run_battery.py`, whose canonical 219/219 record and log are committed at `logs/test_battery.log` and whose v4.3 extension adds the sweep-coverage checks of `tests/test_lag_sweep_artifact.py`), and smoke-validated. Five of them were additionally *executed* end-to-end in the independent re-execution of Section 5.6 and their results are reported in Section 6: (ii) the validation-frozen fixed-FPR operating-point protocol of Section 4.8 (realised test behaviour in Section 6.13); (iii) the untouched-holdout client evaluation mode, in which every client reserves a fraction its models never train on, so local versus global comparisons are clean generalisation estimates (Section 6.16); (iv) the budget-matched centralised retraining (epoch cap raised to 500 in two early-stopping variants; Section 5.5); (vi) the calibration-comparison harness (raw, prior-shift, Platt, isotonic, and temperature arms, validation-fit and single-pass test evaluation; Section 6.13); and the $\alpha=5.0$ promotion run at the full 50-round protocol (negative result; retained $\alpha=1.0$; Section 5.3). The remaining extensions are no longer blocked: (v) the event-level evaluation module (detection rate, missed events, duplicate alerts, warning lead time, alert rate per day) was unblocked by the raw-metadata audit of Section 5.7—event keys and GOES peak times recovered from the raw benchmark—and has been executed on the trained models of both substrates (Tables 5 and 7); and (i) the region-disjoint train/validation split generator is released and awaits a within-fold study design (the temporally-preceding fold is already region-disjoint by construction, Section 5.7, so its remaining use is within-fold validation work).

5.2 Baselines

Four centralised references span the difficulty range of the task. **Climatology** predicts the training prevalence for every window —the skill floor. **Logistic regression** uses balanced class weights and 2,000 maximum iterations. **Centralised MLP** is the architecture-controlled reference: the identical SolarMLP (29,377 parameters), trained on the pooled features with the same loss family and decision layer — the optimisation itself differs as Table 3 records (constant-rate AdamW with early stopping versus cosine-warm-restart rounds; a v4.6 correction of the cell that previously claimed the schedule for both sides) — and its comparison against the federated MLP isolates the effect of federation from the effect of architecture (the v2.x manuscript lacked this baseline and compared federated MLPs only against XGBoost). **XGBoost** [31] (300

Table 2: Frozen configuration of the improved framework (`improvements` branch).

Component	Setting	Notes
Federation	$K = 6$ clients (all selected each round), $T = 50$ rounds	headline
Local optimisation	10 local epochs, AdamW, lr 5×10^{-4} , wd 10^{-4} , cosine warm restarts, batch 512, grad-clip 5.0	unchanged
FedProx	$\mu = 0.01$	sweep-validated
Aggregation	plain size-weighted mean (DA-FL: ablation)	protocol hardening
Loss	Fed-Focal, $\gamma = 2.0$, α clamped (0.05, 0.25); global prevalence computed, not hardcoded	B13 fix
Local rebalancing	SMOTE off (inert: pre-balanced data)	ablation
Partitioning	Dirichlet $\alpha = 1.0$, disjoint by construction, seed 42	B2 fix
MLP	$144 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 1$, BatchNorm, ReLU, dropout 0.3, 29,377 params	
Checkpointing	restore best validation F1 round	B3 fix
Calibration	prior-shift (validation-fit; Platt/isotonic/temperature available)	
Decision layer	F_2 threshold grid $\tau \in [0.05, 0.95]$ on <i>validation</i>	B1 fix
Multi-seed	seeds 42–46, mean/std/CI95, Wilcoxon paired	new
Ablations	$2 \times 3 \times 2 \times 2$ grid, 20 rounds	new

estimators, depth 6, learning rate 0.05, subsample 0.8) is the strongest reference baseline; its `scale_pos_weight` (1.05) is now derived from the *training* prevalence rather than hand-set with knowledge of the test imbalance. Every baseline receives the same validation-selected calibration and threshold treatment as the federated models, so the decision layer is uniform across all comparisons.

5.3 Federated configuration and protocol freeze

Table 2 summarises the frozen configuration. The headline runs use plain size-weighted FedAvg aggregation (the DA-FL weighting of Eq. 7 is studied as an ablation arm in Section 6.15); per-client SMOTE is disabled because the balanced benchmark renders it inert. The proximal coefficient is $\mu = 0.01$. Two facts govern the heterogeneity parameter: the headline run and *all* statistical studies (multiseed, ablations, client-level evaluation) use the frozen-protocol Dirichlet $\alpha = 1.0$; separately, the μ/α sensitivity sweep ($\mu \in \{0.001, 0.01, 0.1\}$, $\alpha \in \{0.5, 1.0, 5.0\}$, 15 rounds) selected $\alpha = 5.0$, $\mu = 0.01$ as the validation-optimal configuration (validation PR-AUC 0.993; its single test pass gave PR-AUC 0.297, consistent with the headline FedProx result). The two values answer different questions— $\alpha = 1.0$ is the adversarially heterogeneous *study regime*, $\alpha = 5.0$ is the validation-selected *tuning optimum*—and we do not silently substitute one for the other. The promotion experiment was executed in this revision at the full 50-round frozen protocol (`outputs/alpha5_promotion.json`): FedProx at $\alpha = 5.0$ reaches test ROC-AUC 0.955 / PR-AUC 0.296, statistically indistinguishable from the $\alpha = 1.0$ headline (0.954/0.307), so the sweep optimum does not transfer to the study regime and $\alpha = 1.0$ is retained as the frozen headline configuration. The same run isolates a severity threshold for the shipped-eval FedAvg drift at $\alpha = 1.0$: at $\alpha = 5.0$ the identical FedAvg arm completes all 50 rounds without the evaluated collapse (test ROC-AUC 0.915, PR-AUC 0.201)—milder heterogeneity keeps the weights compatible with the stale-buffer evaluation for longer, which is consistent with (though not proof of) the mechanism dissected in Section 6.11. Both values, and the provenance of each, are recorded in the run manifest.

Table 3: Training-budget audit of the two architecture-matched arms (review finding: the comparison is confounded by optimisation budget; the table makes the confound auditable).

	Centralised MLP	Federated (FedProx)
Epochs / rounds	≤ 30 epochs	50 rounds $\times$ 10 local epochs
Batch size	256	512
Learning rate (peak)	5×10^{-4}	5×10^{-4}
LR schedule	constant (no scheduler)	cosine warm restarts
Effective passes over own data	≤ 30 (pooled set)	500 (per client shard)
Early stopping	patience 5 (validation BCE)	none (fixed rounds)
Checkpoint rule	best validation BCE	best validation F1 round
Budgets matched	Shipped: no; matched 500-epoch re-run executed—see text	

5.4 Metrics

We report precision, recall, F1, F_2 , ROC-AUC, PR-AUC, Brier score, expected calibration error (ECE), and recall at fixed false-positive budgets (0.5%, 1%, 2%, 5%). Given the $26\times$ validation-to-test prevalence shift and the threshold-transfer failure documented in Section 6.13, the threshold-independent measures (ROC-AUC, PR-AUC, recall-at-FPR) carry the discrimination claims, while the threshold-conditional measures document the decision layer honestly. Confusion matrices are row-normalised. For the multi-seed study we report mean $\pm$ standard deviation with Student- t 95% confidence intervals and one-sided Wilcoxon signed-rank paired tests.

5.5 Optimisation budgets: centralized versus federated

The centralised-MLP versus federated comparison is only interpretable if the optimisation budgets of the two arms are stated explicitly, so the pipeline now emits a machine-readable budget audit (`training_budget` in `results.json`). Table 3 summarises it. The shipped budgets are *not* matched: the centralised reference executes at most 30 full passes over the pooled training set (with early stopping on validation loss), whereas each federated client’s shard receives $50 \times 10 = 500$ local passes over the 50 rounds. The budget-matched re-run was executed in this revision (`outputs/budget_matched.json`): the centralised MLP was retrained on the same pooled data with the epoch cap raised to 500, in two variants—(a) early stopping retained, which fires at epoch 13 and reproduces the headline numbers exactly, and (b) the FL-style fixed-budget protocol (no early stop, best-validation-loss checkpoint retained), which reaches test ROC-AUC 0.917 / PR-AUC 0.219. Two conclusions follow. First, the confound ran *against* the federated arm: the shipped centralised reference was *undertrained* (early stopping fires long before the budget binds), and the pooled model improves when granted the federated arm’s budget. Second, the improvement is not sufficient—at full budget parity the centralised MLP (0.917/0.219) remains below FedProx (0.954/0.307) on the threshold-free metrics, so the headline gap is a genuine federation effect on this benchmark rather than a budget artifact (its F_1 of 0.135 versus FedProx’s 0.042 reflects threshold-transfer behaviour under prior shift, not ranking skill). The architecture comparison (identical 29,377-parameter MLP) remains matched by construction. The raw-substrate LSTM arms (Section 6.5) inherit the same asymmetry: the pooled SolarLSTM comparator is capped at 30 epochs with patience-5 early stopping (which did not fire: epoch 30, validation BCE 0.039), while the federated arms receive 500 shard passes. The FedProx-LSTM-over-central-LSTM margin (+0.008 ROC-AUC) is therefore budget-caveated—the cleaned-fold parity experiment moved the central MLP by +0.029 ROC-AUC, the same order—whereas the load-bearing contrast of that section (federated LSTM versus federated MLP under identical budget and aggregation) is budget-matched by construction.

5.6 Compute and reproducibility

The complete experimental programme of this revision (headline run, 24-cell ablation grid, 9-configuration sweep, 5-seed study, interpretability, and communication measurements) was executed on a 2-core CPU-only container with under 4 GB of RAM, in aggregate approximately 2.5 hours of wall-clock time. Every experiment writes a machine-readable results artefact (`results.json`, `ablation_results.json`, `sweep_results.json`, `multiseed_results.json`, `interpretability.json`, `smote_validity.json`) plus a run manifest embedding the code version, seeds, configuration, and dataset manifest hash; all figures and tables are generated from these artefacts. The pipeline supports round-level crash recovery, and the integrity battery (partitioning disjointness, leakage audit, calibration, evaluation protocol, threshold-freezing arithmetic, event-level metrics, FL smoke tests, GPU-queue artefact guards, a leakage-gate failure-path test, and an audit-artefact recomputation test, joined at v4.3 by a lag-definition-sweep cross-artefact test) passes its canonical 219/219 record (`tests/run_battery.py`; `logs/test_battery.log`, torch-equipped) and 224/0 in the v4.3 torch-less verification environment (`logs/test_battery_v4.3.log`, 24 torch-dependent checks skipped — the same delta an independent torch-less re-run shows).

Independent re-execution from the public artefacts. Following an external review, the entire headline pipeline was re-executed from scratch in a clean environment: the Cleaned SWAN-SF release was downloaded from the dataset repository’s official distribution, all 20 partition files were verified byte-identical to the frozen data manifest (SHA-256 match on every file — `data_manifest/verify_manifest.py`, whose committed log at `logs/verify_manifest.log` records the 20/20 verdict and the manifest hash; the verification also runs as a hard gate at pipeline start), and the frozen protocol was run end-to-end (data load and flattening, immutable splits, Dirichlet partition, pooled baselines, FedAvg and FedProx to 50 rounds, validation-fit calibration and thresholds, single-pass test evaluation, figures, and SHAP attribution) on CPU only (the protocol run of record used Python 3.12; the committed battery log header records the v4.1 verification battery’s environment — python 3.13.5, numpy 2.2.4, torch 2.14.1+cpu — the library versions pinned in `requirements.txt`; the run manifest records them per-run from v4.0 onward — the committed manifest documents the original v3.x execution and predates the field). Every test-set metric of every model reproduced *exactly*, to all reported significant digits: all six models’ ROC-AUC, PR-AUC, Brier, ECE, precision, recall, F1, and accuracy match the published artefacts at three decimals, and the FedAvg divergence signature reproduces (rounds 5–40 agree to three decimals; the collapsed-tail values differ by up to 0.02 between the original execution and the re-run caches, a chaotic regime the diagnostic of Section 6.11 documents). The pipeline is therefore bit-stable under this protocol: fixed seeds, deterministic CPU kernels, disjoint index arithmetic, and single-precision training make the published numbers re-derivable by anyone from the public dataset and the released code, which is the strongest form of reproducibility claim this study can make. The same re-execution session ran, for the first time on real data, the calibration-arm comparison (Table 16), the untouched-holdout client evaluation (Section 6.16, `client_holdout_eval.json`), and the frozen-FPR operating points now reported in Section 6.13; the multi-seed, sweep, and ablation studies were not re-executed and their numbers remain the owner-side artefacts embedded in the repository.

Owner-GPU execution of the sequence arms. The raw-substrate LSTM arms of Section 6.5 were executed on owner-side GPU hardware (an NVIDIA RTX 4060 Laptop with CUDA and cuDNN), because the CPU-only protocol above would have been impractically slow for the sequence arms—the one experiment where GPU was a practical requirement. Two batches: the seed-42 arms (2.6 h total: 0.08 h pooled comparator, 1.16 h FedAvg-LSTM, 1.33 h FedProx-LSTM) and, as a second single resumable queue, the remaining programme (8.7 h: 1.5 h

SCAFFOLD-LSTM, 3.9 h the seed-43 replication of all four arms, 3.2 h the two SMOTE ablation arms). The substrate identity is device-independent and asserted at startup on every run (label equality with the frozen 2D split; a 144-statistic feature cross-check that runs whenever the 2D cache is present, asserting $\max|\text{diff}|$ below the 5×10^{-3} float16 bound — on the owner-GPU box only the 3D cache was shipped, so the label-equality assertion ran there and the feature check ran in the CPU sessions), and each run writes the same machine-readable artefact schema as every CPU arm (`raw_lstm_eval`, `raw_lstm_scaffold`, `raw_lstm_seed43`, and `raw_lstm_smote`, all under `outputs/`). Each GPU batch was validated on receipt, before integration, by a machine artefact audit (64 checks for the second batch: protocol fingerprints, substrate identity, embedded frozen-counterpart bit-matches, event-key reconciliation, and cross-checks against the run logs). The residual device effect is kernel-level nondeterminism in cuDNN recurrence and float-reduction order, predicted within ± 0.001 – 0.005 on AUC-scale metrics and now measured by the seed-43 replication at ≤ 0.005 ROC-AUC on every arm (Section 6.6); all other artefacts in this paper were produced under the bit-stable CPU protocol above.

BatchNorm-transport diagnostic and no-BatchNorm control. Two CPU-only experiments extend the frozen protocol without altering it. The diagnostic runner (`experiments/run_bn_diagnostic`) loads the shipped checkpoints, audits their BN buffer state, and evaluates the returned artefacts and a faithfully re-run FedAvg trajectory (validation history matches the frozen run within ± 0.01 ROC-AUC at every logged round) under four BN treatments: init buffers (the shipped path), exact per-client statistics pooled by the law of total variance, per-batch statistics at evaluation time, and the element-wise weighted average of the per-client running statistics that local training produced in the round being evaluated. Recalibration uses training shards only; validation and test are never touched by it. The control runner (`experiments/run_nobn_control.py`) replaces the three `BatchNorm1d` layers with identity (28,929 parameters vs. 29,377) and runs the identical 50-round loop for both FedAvg and FedProx, with the shipped best-validation checkpoint rule as the only selection mechanism; its per-round test evaluation is a disclosed diagnostic, not a selection path. Both runners are round-level resumable and write machine-readable artefacts (`bn_diagnostic.json`, `nobn_control.json`) under `outputs/`.

5.7 Data-provenance audit and the leakage-free protocol

Raw-benchmark alignment. The Cleaned SWAN-SF export carries no region or event identifiers, so the review-programme items that require them (region-disjoint splits, event-level evaluation, natural-prevalence substrates) were queued pending raw metadata. We obtained the raw benchmark from its official distribution (Harvard Dataverse, doi:10.7910/DVN/EBCFKM, 6.5 GB across the five partition archives plus the 34 MB addenda with the GOES flare event lists) and aligned every cleaned window to its raw instance. Each raw instance is a 60-timestep CSV whose filename encodes the flare class, the *HARP* region identifier, and the window start/end timestamps, and the folder (FL for M/X-flaring, NF otherwise) gives the label. The matching key is the per-feature argmax/argmin position along the 60 timesteps—48 discrete invariants that survive any per-feature monotone normalization and are robust to the export’s imputation noise. A window is accepted as aligned when at least 28 of 48 invariants agree (borderline cases require a mean per-feature Spearman rank correlation of at least 0.70). On the test partitions this verifies 98.7–100.0% of windows with 100.00% label agreement (train-side label agreement is 99.71% — 161 disagreements, which bound the synthetic-share estimate below) and a median best-to-second-best margin of 31–32 invariants; the alignment artefacts and per-window metadata tables are released with the code (`outputs/dataset_structure_audit.json`; a committed test, `tests/test_audit_artifact.py`, recomputes every headline audit total from the committed per-partition tables and slim metadata).

What the audit found. Three structural facts follow. First, the test export of partition p contains *all* raw instances of p (the counts match exactly on every partition). Second, the training export of p is a RUS-Tomek-TimeGAN rebalanced *subset of the same instances*: the verified training pool contains 56,181 windows covering 56,005 unique raw instances (the 176-window difference is duplicate-instance windows, 160 of them positive, all in partition 4), and *every* unique verified train instance is also a test instance (56,005/56,005; on the row basis, 56,181/56,181). Every M/X-flaring test instance—6,234 of 6,234, 100% of the minority class—is in the training pool. Third, 85.7–90.0% of the training positive windows are TimeGAN-synthetic (they have no raw counterpart; the 161 train-side label disagreements bound this estimate), while virtually no negative windows are synthetic. The release paper of the cleaned dataset prescribes pairing *temporally preceding* training partitions with a later test partition; the shipped pipeline of this study instead pooled the training exports of all five partitions and evaluated on the test exports of the same five, pairing each test partition with a rebalanced subset of itself. The published in-partition numbers therefore partially measure memorised data, most severely for the minority class. Because no region identifier spans two partitions (verified: 0 shared HARPs), a temporally-preceding fold is automatically instance-, region-, and time-disjoint.

Leakage-free fold. All headline-model comparisons were re-run under the benchmark’s intended protocol: training on the training exports of partitions 1–4 (77,865 windows; 65,406 train / 12,459 validation after the identical stratified carve), and a single-pass evaluation on the test export of partition 5 (75,365 windows at the natural 1.31% prevalence). Everything else is frozen: Dirichlet $\alpha = 1.0$, six clients, 50 rounds, seed 42, Fed-Focal loss, prior-shift calibration and validation-frozen thresholds. The runner is `experiments/run_partition_disjoint.py` (round-resumable); its artefacts are `outputs/partition_disjoint_eval.json` and, for the event-level analysis, `outputs/event_level_p5.json` with the per-window event keys recovered from the raw filenames and the addenda GOES peak list (65 of 65 partition-5 flare events matched, 4 with multiple candidates). Event grouping uses the filename flare tag for positive windows (one group per physical flare) and region-day groups for background windows; alert de-duplication applies a 1-hour cooldown within each group.

6 Results and Analysis

All numbers in this section are produced by the improved pipeline (improvements branch) under the frozen evaluation protocol of Section 5: a single immutable train/validation/test split, thresholds and calibration selected on validation only, and the test set touched exactly once per model. Every table and figure is generated from the machine-readable results-json artefacts of the runs; no value is transcribed by hand. The headline run uses 50 communication rounds, six clients, and seed 42; Sections 6.14 and 6.15 quantify seed sensitivity and component attribution. An independent re-execution of the headline pipeline from the public, byte-verified dataset artefacts reproduces every test metric of every model exactly (Section 5.6); the calibration, frozen-FPR, and client-holdout results below were produced in that re-execution session.

Protocol disclosure. The data-provenance audit of Section 5.7 established that the in-partition train/test pairing used by the shipped headline protocol shares instances between the training pool and the test set—including 100% of the flaring test instances. The subsections that follow the present one therefore report *in-partition* results: they remain valid as evidence of protocol stability, seed sensitivity, component attribution, and calibration behaviour under the frozen pipeline, but they are not estimates of generalisation to unseen data. The leakage-free evaluation under the benchmark’s intended temporally-preceding protocol is reported first, and the claims of this study are tiered against it in Section 6.7.

Table 4: Leakage-free fold (train partitions 1–4, single-pass test on partition 5, natural 1.31% prevalence) against the shipped in-partition numbers. Threshold-conditional calibration metrics (Brier, ECE) use the validation-frozen prior-shift pipeline; ROC/PR are threshold-free.

Model	In-partition (shipped)		Leakage-free fold				
	ROC	PR	ROC	PR	Brier	ECE	R@2%FPR
Logistic regression	0.824	0.235	0.978	0.455	0.375	0.477	0.695
XGBoost (centralised)	0.977	0.493	0.971	0.457	0.055	0.101	0.692
Centralised MLP	0.888	0.153	0.973	0.382	0.544	0.687	0.555
FedAvg MLP	0.575	0.199	0.906	0.370	0.127	0.337	0.652
FedProx MLP	0.954	0.307	0.976	0.308	0.271	0.508	0.631

6.1 Leakage-free evaluation on the temporally-preceding fold

Table 4 compares the five headline arms under the leakage-free fold (training partitions 1–4, single-pass test on partition 5) against the shipped in-partition numbers. Three observations follow. First, *discrimination generalises*: every arm attains a ROC-AUC of at least 0.906 on the unseen partition, and the task is evidently easier than the in-partition protocol suggested (logistic regression reaches 0.978 / PR-AUC 0.455, the best PR ranking of all arms together with XGBoost’s 0.457). Second, the centralised-versus-federated gap of the in-partition protocol *disappears*: the pooled centralised MLP of identical architecture reaches 0.973 / 0.382 against FedProx’s 0.976 / 0.308—a ROC difference of 0.003 and a reversed PR ordering. The in-partition gap (0.888 vs 0.954) was an artifact of the protocol pairing, not a federation effect; the budget-parity re-run of Section 5.5, which controls optimisation budget but not instance overlap, must be read in the same light. Third, the FedProx-versus-FedAvg contrast partially survives: FedProx’s ROC-AUC (0.976) exceeds FedAvg’s (0.906) by 0.070, but FedAvg’s PR-AUC (0.370) is *higher* than FedProx’s (0.308). The nuance is in the checkpoint dynamics: plain FedAvg’s validation F1 collapsed to zero by round 40 and the protocol restored its round-35 checkpoint (validation F1 0.915), so its leakage-free outcome depends on which round the validation monitor happened to freeze. We no longer describe either arm’s outcome as checkpoint-independent: the BatchNorm diagnostic of Section 6.11 shows the validation F1 monitor saturates at the all-positive value $2p/(1+p)$ whenever the evaluated probabilities saturate, and the checkpoint rule then acts on a saturated signal—FedProx’s apparent safety on this fold is that its evaluated probabilities never crossed that regime, not that its artefact is selection-invariant.

The calibration and threshold-transfer failures documented in-partition (Section 6.13) persist leakage-free — on this fold the balanced-validation curve (48.9% prevalence, partitions 1–4) again faces the natural-prevalence partition-5 test (1.31%), the same 26-fold prior shift without the instance overlap — and are therefore genuine properties of the balanced-training / natural-test interface rather than overlap artifacts: the validation-frozen FPR operating points are violated by every arm (e.g. XGBoost realises 15.3% FPR at the frozen 2% target; FedProx 49.4%), and only XGBoost’s posteriors remain usable (Brier 0.055, ECE 0.101 versus 0.127–0.544 for the neural arms; FedAvg’s 0.127 sits closest but is still an order of magnitude above the 0.013 climatology floor). The prior-shift-corrected neural posteriors are not repairable by threshold re-freezing on this fold either.

Standard verification metrics and baselines. Because the toolkit of Leka et al. [6]—TSS, HSS, Brier skill, reliability, persistence and climatology—is the standard vocabulary of flare-forecast evaluation, it accompanies every stored operating point of every evaluation substrate (the F_β thresholds and frozen-FPR points of the in-partition, leakage-free and raw-substrate 2-D arms, the ten raw-substrate LSTM cells, and the event-level tables; 10-bin reliability is recomputed for the six in-partition arms from the phase-cache models;

`experiments/run_standard_metrics.py` → `outputs/standard_metrics.json`, strict JSON; the artefact-derived blocks (A, C, D) are content-exact on re-execution, block B’s reliability recomputation additionally requires the uncommitted phase-cache models, and the artefact embeds a wall-clock `elapsed_s` field, so file-level equality is not claimed). Four findings matter. First, on the leakage-free fold at each arm’s frozen F_β threshold, XGBoost attains TSS 0.848 / HSS 0.201 (best frozen-FPR point 0.847), logistic regression TSS 0.489 / HSS 0.025, while the neural arms sit at TSS 0.06–0.19 / HSS ≤ 0.006 : the prevalence-free TSS ranking matches the PR-AUC ranking at the top of the table (XGBoost, logistic regression, then the neural arms), but HSS prices the neural arms’ false-alarm load severely at 1.31% prevalence. Second, the persistence baselines—computed on the true time order within each active region, as the operational verification culture mandates—set two floors, and neither is decorative. Same-region previous-window persistence (label inertia) achieves TSS 0.967 / HSS 0.970 on the partition-5 test with no learned model at all: at a 24-hour label horizon on ~ 60 -minute windows, an active region that is flaring is still flaring, and *no arm in this paper—XGBoost included—beats that floor*. The 24-hour-lagged variant (forecast = the label of the same-region window nearest $t - 24$ h, matched within ± 35 minutes) scores TSS 0.418 / HSS 0.434; on TSS, XGBoost (0.848) and logistic regression (0.489) are the only arms above it, and every neural arm at its frozen threshold sits below it, while on HSS no arm clears the floor at all (XGBoost’s HSS is 0.201). This is the classic Leka-era result, and a benchmark audit is the right vehicle for stating it plainly: at this window geometry the honest question is not whether models beat persistence, but whether anything beats *lagged* persistence—and on this fold, at TSS, two arms do, none at HSS. (An earlier implementation of this baseline ordered windows by pool row, whose within-region order is temporally scrambled—the median gap between pool-row-adjacent windows is two days—so its stored value measured randomly paired windows; that is precisely the class of evaluation-protocol defect this paper exists to audit, and the corrected ordering and both variants regenerate from the committed artefact. A sensitivity sweep over eight candidate lag definitions is committed alongside it (`experiments/run_lag_definition_sweep.py` → `outputs/lag_definition_sweep.json`): the seven point-lag variants land at TSS 0.392–0.418 with the same two arms clearing each and none at HSS, while a lookback-flare variant scores 0.964—label inertia under another name, which no arm clears—so the paper reports inertia and lagged persistence as two distinct floors.) Third, no arm’s raw Brier beats the climatology floor (pooled-test BSS: XGBoost -2.4 , FedAvg -6.3 , FedProx -13.3 ; all negative), and the 10-bin reliability tables confirm the systematic over-confidence that the calibration analysis of Section 6.13 dissects. Fourth, event-level detection rates carry Wilson 95% half-widths of ± 2.8 – 4.0 points on the leakage-free fold and up to ± 11.8 points on the raw substrate (65 events), and alert rates carry exact chi-square (Garwood) Poisson intervals, so the between-arm event-level narrations of the next paragraph are read against those bands.

Event-level evaluation. Table 5 regroups the partition-5 test decisions into physical flare events using the raw-benchmark event keys (65 M/X-class events; background windows grouped by region and day; 1-hour alert cooldown). All five arms detect effectively every event (64–65 of 65) with a median first-alert lead time of 23.4 h before flare peak. That lead must be read for what it is: the label is positive for *every* window in the 24 hours before a flare, so a median first alert near the horizon’s edge measures the label boundary, not precursor physics — an honest reading is “detection fires early in the labelled window,” not “the model sees the flare 23 hours ahead.” What separates the arms is their false-alarm burden: XGBoost emits 1.7 false-alarm windows per day versus 11.6–12.7 for the neural arms, a $7\times$ difference that mirrors the window-level calibration gap and confirms that the deployability boundary on this benchmark is calibration, not detection. After the 1-hour cooldown, roughly 70% of retained alerts are within-event duplicates, quantifying the persistence behaviour an operations desk would face.

Table 5: Event-level metrics on the leakage-free fold (65 true M/X-class events, validation-frozen F_2 thresholds, 1-hour alert cooldown). FA = false alarm.

Model	Events detected	FA windows/day	Alerts/day	Lead to peak (h)
Logistic regression	65/65 (100%)	8.6	8.8	23.4
XGBoost (centralised)	64/65 (98.5%)	1.7	1.9	23.3
Centralised MLP	65/65 (100%)	12.7	12.8	23.4
FedAvg MLP	65/65 (100%)	11.6	11.8	23.4
FedProx MLP	65/65 (100%)	11.9	12.1	23.4

6.2 Raw-substrate retraining: the unbalanced benchmark

The leakage-free fold of Section 6.1 still trains on the cleaned export—RUS–Tomek rebalanced, TimeGAN-augmented, pre-normalised. The raw-benchmark retraining (re-run item 9, executed in this revision) asks the complementary question: does anything in the comparison survive when the frozen protocol is trained directly on the *raw* SWAN-SF partitions at their natural class prevalence (2.05% across training partitions 1–4), with no rebalancing, no synthetic positives, and the two cleaning stages reproduced in-pipeline? The raw benchmark cannot be trained on directly: 8.87% of its values are missing (42.3M of 476.9M; TOTFZ is unobserved in 98.7% of training windows) and the 24 parameters span 11 orders of magnitude. The experiment therefore reproduces the release’s two preprocessing stages from their published description (ApJS, 10.3847/1538-4365/ad7c4a): a *fast Pearson-correlation-based k-NN imputation* (partial gaps filled by within-window interpolation; fully-missing series transferred from the $k=5$ nearest training windows in the space of the five most correlated observable features) and the *LSBZM* normalisation chain (robust clip, shift, log/sqrt/Box–Cox dispatched by fitted λ , z-score, min–max), with every parameter fitted on the training pool only. A verification against the released cleaned export on the audit-matched partition-5 windows (100% window match) bounds the reproduction: at positions where the raw value was observed and nonzero, rank agreement with the release has median Spearman $\rho = 0.898$ (min 0.842; the release is itself an exact monotone image of the raw values, $\rho = 1.000$, so the residual gap is float16-storage tie dilution); at positions the release imputed, agreement is method-dependent (median $\rho = 0.313$), and for the R-value flux-emergence proxy the release re-imputes the entire zero mass (60.9% of that column) as if it were missing, whereas this reproduction preserves zeros as physical values. The three substrates therefore differ in composition as well as in preprocessing provenance, which is precisely the contrast the experiment isolates.

Table 6 reports the six arms under the identical frozen protocol (Dirichlet $\alpha=1.0$, six clients, 50 rounds, Fed-Focal, seed 42, prior-shift calibration, validation-frozen thresholds). The result is decisive and asymmetric. *Centralised arms are substrate-robust*: logistic regression is unchanged to the third decimal (0.978, PR 0.448 vs 0.455), the centralised MLP retains 0.971 ROC-AUC and its PR-AUC *rises* from 0.382 to 0.445, and XGBoost keeps 0.974 ROC-AUC while its PR-AUC falls from 0.457 to 0.329—the synthetic TimeGAN positives were inflating its minority-class precision, not teaching it. *Federated arms are substrate-fragile*: FedProx degrades from 0.976/0.308 to 0.930/0.149; plain FedAvg’s PR-AUC collapses from 0.370 to 0.056 (its calibrated posteriors never cross the validation-frozen threshold: test recall 0.000); and SCAFFOLD—the one re-run item that had never been executed on any real substrate—fails outright at 0.768/0.057 with Brier 0.137 (every SCAFFOLD cell in this section measures the implementation’s five disclosed departures from the reference algorithm, Section 4, and is read below as implementation-case evidence). The federated-versus-centralised ordering of the leakage-free fold was still a property of the balanced training substrate; on the raw unbalanced benchmark the focal-loss clients cannot learn the 2% minority class through 6-client Dirichlet heterogeneity, and the gap between the best federated and best centralised arm widens to 0.048 ROC-AUC and $3\times$ PR-AUC.

Table 6: Raw-substrate retraining (train raw partitions 1–4, natural 2.05% prevalence, FPCKNN/LSBZM reproduced in-pipeline with train-only parameters; single-pass test on raw partition 5, natural 1.31% prevalence; identical frozen protocol). The leakage-free cleaned-substrate fold is reproduced for comparison; SCAFFOLD was not run on the cleaned fold.

Model	Fold (cleaned)		Raw substrate			
	ROC	PR	ROC	PR	Brier	R@2%FPR
Logistic regression	0.978	0.455	0.978	0.448	0.053	0.732
XGBoost (centralised)	0.971	0.457	0.974	0.329	0.018	0.587
Centralised MLP	0.973	0.382	0.971	0.445	0.010	0.596
FedAvg MLP	0.906	0.370	0.875	0.056	0.013	0.063
FedProx MLP	0.976	0.308	0.930	0.149	0.032	0.343
SCAFFOLD MLP	—	—	0.768	0.057	0.137	0.175

One disclosure this section carried from v4.8, now resolved: the three federated *MLP* rows of Table 6 are evaluated through the same untransported-BatchNorm path that Section 6.11 indicts in-partition — the runner transports model *parameters* only, so the returned checkpoints carry never-updated running buffers into a `model.eval()` evaluation. Until v4.9 the raw-substrate MLP degradation above was therefore *BN-evaluation-conditional* on its face; the checkpoint-level diagnostic has since been executed (Section 6.3, `raw_bn_diagnostic.json`) and resolves the evaluation-path question *against* the BN explanation: no evaluation-time treatment recovers the centralised reference. The deficit itself, however, is neither evaluation-conditional nor encoder-conditional — the no-BatchNorm re-training of Section 6.4 removes the three BN layers and nothing else, and the federated MLP recovers to the sequence arms’ ranking. The sequence arms (SolarLSTM carries no BatchNorm) were never exposed to the mechanism, which is exactly why they never collapsed on this substrate.

The event-level evaluation on the same substrate (Table 7) makes the collapse operational: all 65 true M/X-class events are matched to their GOES peaks exactly as in the cleaned fold, but plain FedAvg now detects *zero* of them, FedProx detects 48 (73.8%) at 1.3 false-alarm windows per day, and SCAFFOLD’s 83.1% detection rate is bought at 6.4 false-alarm windows per day—roughly 20× the centralised arms’ burden. The centralised arms detect 61.5–80% of events at 0.2–0.3 false-alarm windows per day with 18–23 h median lead times. The sequence arms appended below the rule (Section 6.5) restore partial federated detection and are discussed there. Detection under class balance, in other words, was the artifact; under the benchmark’s natural prevalence the deployable models are exactly the ones that never needed federation.

6.3 Raw-substrate BN diagnostic: resolving the v4.8 condition

The v4.8 revision shipped the runner for this diagnostic (`experiments/run_raw_bn_diagnostic.py`) and queued its execution owner-side; it has since run (RUNLOG ask #4), and its artefact of record — `outputs/raw_bn_diagnostic.json`, sha-pinned by the battery, together with the retrain record it reproduces, `outputs/raw_substrate_rerun.json`, pinned under the same convention — is what this subsection reads. One structural fact precedes the numbers, disclosed at equal length in the RUNLOG: the v3.x-era raw-substrate checkpoints were *never persisted* (that era’s runner stored the 3D caches and the sequence-arm checkpoints, not the seed-42 MLP weights), so the same-weights counterfactual the diagnostic was designed around is unavailable by construction. The run therefore re-evaluated a fresh seed-42 replication: the substrate was rebuilt from the surviving 3D cache (one additional f16 quantisation, within the 5×10^{-3} series bound the cache format documents), the three arms were retrained round-resumably under the frozen protocol, and the diagnostic then applied its three BN treatments to those identical weights. The replication is close at the substrate level: the pooled baselines reproduce

Table 7: Event-level metrics on the raw substrate (65 true M/X-class events, validation-frozen F_2 thresholds, 1-hour alert cooldown; 99.86% of test windows matched to the audit’s event keys). The sequence arms (below the rule) are evaluated by the standalone torch-free CPU runner over the owner-GPU run’s stored per-window test probabilities and the same audit metadata; the SCAFFOLD-LSTM row was evaluated in-runner by the same event-level machinery during the second GPU batch. Median leads are rounded to 0.1 h except the SCAFFOLD-MLP cell, whose stored median (1,395 min) is exactly 23.25 h and is shown at full precision (T2 erratum, v4.8); the FedAvg arm-C row (batch-statistics evaluation, Section 6.3) is the companion cell to its shipped zero-detection row.

Model	Events detected	FA windows/day	Alerts/day	Lead to peak (h)
Logistic regression	51/65 (78.5%)	0.3	0.4	22.8
XGBoost (centralised)	40/65 (61.5%)	0.2	0.3	18.0
Centralised MLP	52/65 (80.0%)	0.3	0.4	20.5
FedAvg MLP	0/65 (0.0%)	0.0	0.0	—
FedAvg MLP (arm C, batch statistics)	38/65 (58.5%)	0.28	0.35	19.8
FedProx MLP	48/65 (73.8%)	1.3	1.4	23.0
SCAFFOLD MLP	54/65 (83.1%)	6.4	6.6	23.25
Centralised LSTM	41/65 (63.1%)	0.3	0.4	19.3
FedAvg LSTM	35/65 (53.8%)	0.2	0.2	17.4
FedProx LSTM	49/65 (75.4%)	0.4	0.5	23.2
SCAFFOLD LSTM	55/65 (84.6%)	0.4	0.5	23.2

Table 8: The raw-substrate BN diagnostic (`raw_bn_diagnostic.json`, v4.9): test ROC-AUC / PR-AUC of the identical replication weights under the three BN treatments, per federated MLP arm. A: the published init-buffer evaluation path; B: exact per-client statistics pooled by the law of total variance (training shards only); C: evaluation-time batch statistics. Gate Δ is arm A’s ROC-AUC delta against the published rows of Table 6 (the disclosed MISMATCH; FedAvg carries the replication’s retrain variance). The SCAFFOLD row measures the implementation’s sign-inverted, cold-start variant (the five departures of Section 4) — an implementation case, not the reference algorithm.

Arm	A: init buffers	B: pooled	C: batch stats	Gate Δ
FedAvg	0.898 / 0.070	0.501 / 0.015	0.915 / 0.227	+0.023
FedProx	0.928 / 0.149	0.500 / 0.013	0.920 / 0.251	−0.002
SCAFFOLD	0.763 / 0.053	0.500 / 0.013	0.757 / 0.069	−0.005

the published rows to $\leq 2.3 \times 10^{-3}$ ROC-AUC (logistic regression $+1.7 \times 10^{-4}$, XGBoost $+1.5 \times 10^{-3}$, centralised MLP -2.3×10^{-3}), which isolates the federated arms’ gate deltas — ROC-AUC +0.023 (FedAvg), −0.002 (FedProx), −0.005 (SCAFFOLD) — as retrain nondeterminism concentrated in the highest-variance arm. The gate recorded MISMATCH on all three arms (tolerance 10^{-4}) and exited non-zero as designed; the record is frozen with that verdict.

Table 8 yields the reading, and it is the opposite of the in-partition pattern of Section 6.11. First, *the raw-substrate MLP deficit is not a manufactured collapse*: arm A — the very path the in-partition diagnostic indicts — holds 0.763–0.928 ROC-AUC on every arm here, where the in-partition collapse scored 0.575, and no treatment approaches the centralised reference (0.969 on the same replication): the residual deficit, roughly 0.04–0.21 ROC-AUC depending on arm and treatment, is a property of the flattened-feature arms under federation, not of the buffer-transport protocol. Second, *pooled recalibration is actively destructive here*: arm B collapses every arm to chance (0.500–0.501), whereas in-partition the same oracle *recovered* FedAvg from 0.575 to 0.893 — pooling exact per-client statistics across 6 heterogeneous raw-prevalence clients inflates variance past the point where any single pooled normalisation ranks the test distribution, so the in-partition “optimistic recalibration rescue” does not transfer to the natural-prevalence substrate at all. Third, evaluation-time batch statistics (arm C) move ROC-AUC by at most 0.017 while sharpening the positive score scale (PR-AUC 0.070→0.227

for FedAvg, 0.149→0.251 for FedProx): a calibration instrument, not a rescue. The diagnostic’s event-level layer also records the operating-point pathology the A-path thresholds inherit — the replication’s checkpoint rule returned best-val-F1 states of 0.0 for FedAvg and FedProx (and since the strict-inequality rule initialises its best at 0.0 and never fires on an all-zero history, the shipped checkpoint is the final round’s global model, round 50, in both cases), FedAvg’s val-frozen threshold sat at the 0.05 floor, and the arm detects 0 of 65 events under A while arm C, the only treatment with a usable validation operating point, detects 38 (58.5%) at 0.28 false-alarm windows per day (FedProx: 62 at 2.8 under A, 50 at 0.43 under C). The v4.8 condition is therefore closed with a precise scope: untransported BN statistics can manufacture collapse — in-partition proved that — but the raw-substrate federated MLP degradation is not that phenomenon, and the sequence arms (no BatchNorm) were never exposed to it.

Two follow-ups complete the reading. First, arm B’s own instrument has been validated: applying the byte-identical pooled-recalibration machinery to the *centralised* raw MLP — where pooled per-client statistics should reconstruct the model’s own running statistics — preserves the centralised ranking to -6.8×10^{-5} ROC-AUC and *improves* PR-AUC by 4.9×10^{-3} (0.96864/0.43051 against arm A’s own-buffer 0.96870/0.42566, both reproducing the retrain record to better than 10^{-6} ; `outputs/arm_b_central_sanity.json`, committed and battery-checked). The supporting diagnostics record the reconstruction’s drift directly: pooled-versus-own buffer deltas run 0.021/0.043/0.113 (mean relative L2) across layers 0–2, depth-monotone and non-trivial at the last layer, and yet the ranking holds — the instrument tolerates realistic statistics drift without collapsing. Arm B’s federated collapse is therefore a genuine property of pooling per-client statistics across six raw-prevalence clients, not an implementation artefact, and the B column of Table 8 is its *validated proxy*: what a transport fix that pooled normalisation statistics would actually deliver on this substrate. The faithful buffer-transport re-run (arm D) remains the endorsed clean-closure option for an external submission, but the reading no longer depends on it. Second, the attribution question this analysis left open — evaluation artefact, encoder, or BatchNorm-under-federation — is answered by the control of the next subsection.

6.4 No-BatchNorm federation on the raw substrate: the attribution control

The diagnostic above fixes the evaluation-path question but leaves the attribution open: the collapsing arms of Table 6 carry BatchNorm, the rescuing sequence arms of Section 6.5 do not, so “the flattened-feature MLP degrades under federation” could not be separated from “BatchNorm degrades under federation” — the FedBN line of work the paper itself cites keeps exactly that alternative alive. The control is one intervention: the three `BatchNorm1d` layers replaced by the identity (the in-partition construction of Section 6.12: same stack minus gamma/beta, 28,929 parameters against the BN baseline’s 29,377, no running statistics, so the shipped evaluation path *is* the exact evaluation), retrained end to end under the frozen raw protocol — identical seed-42 Dirichlet shards, 50 rounds, Fed-Focal, prior-shift calibration, validation-frozen thresholds (`experiments/run_raw_nobn.py`; artefact of record `outputs/raw_nobn_eval.json`, its reference columns read live from the committed artefacts at run time, never hand-typed; both arms executed owner-side, 1,699s total, the FedAvg arm’s 50 rounds preserved through a mid-run crash by the round-resumable state design, and the round-selection bug that caused the crash fixed and battery-guarded — an erratum disclosed in the RUNLOG, not in the numbers).

Removing BatchNorm — and nothing else — recovers the overwhelming majority of the raw MLP’s federation penalty (Table 9): FedAvg gains +0.088 ROC-AUC and +0.310 PR-AUC (6.5×); FedProx gains +0.045 and +0.196 (2.3×). In gap terms: FedAvg’s ROC-AUC gap to the centralised MLP closes from 0.096 to 0.008 (91%) and its PR-AUC gap from 0.389 to 0.079 (80%); FedProx’s ROC-AUC gap closes *past* zero (−0.004 residual, the federated arm above its centralised counterpart) and its PR-AUC gap from 0.296 to 0.100 (66%). Both no-BN arms

Table 9: The attribution control on the raw substrate: test window-level ROC-AUC / PR-AUC under the frozen protocol, seed 42, shipped validation-selected checkpoints. The no-BatchNorm arms are the single-intervention retraining (three `BatchNorm1d` layers $\rightarrow$ identity, nothing else changed); the reference rows are read from the committed artefacts at run time.

Arm (raw substrate)	ROC-AUC	PR-AUC
FedAvg MLP (BatchNorm, published)	0.875	0.056
FedAvg MLP (no BatchNorm, this control)	0.963	0.366
FedProx MLP (BatchNorm, published)	0.930	0.149
FedProx MLP (no BatchNorm, this control)	0.975	0.345
Centralised MLP (reference)	0.971	0.445
FedAvg LSTM (reference)	0.958	0.305
FedProx LSTM (reference)	0.970	0.404

sit at or above their sequence counterparts on ROC-AUC (0.963 vs 0.958; 0.975 vs 0.970), and no-BN FedProx exceeds the *centralised* MLP on ROC-AUC (0.975 vs 0.971). The federated collapse of the BN raw MLP arms is therefore dominated by BatchNorm-under-federation — not by the flattened-feature encoder, not by the evaluation path (Section 6.3), and not by federation per se. The sequence arms’ rescue of Section 6.5 is real but is one instance of the BatchNorm-free rescue, of which this control is the direct proof; SCAFFOLD’s raw-substrate failure (0.768/0.057) is the one federated collapse this control has not been run against, so its reading stays implementation-case-qualified as before.

Checkpoint-selection hygiene is born with these arms and disclosed in full, because the degenerate-monitor pathology of the BN arms (Section 6.3) is what it guards against. The shipped best-val-F1 rule selects round 50 for both arms (FedAvg’s validation F1 0.527; FedProx’s 0.092 — its validation scores are compressed at the 0.35 default threshold, and its calibrated test operating point sits at 0.155, test F1 0.400 at recall 0.763), and a validation-ROC-selected sensitivity row rides beside it (FedAvg: round 40, 0.9635/0.3592; FedProx: round 50, identical to the shipped row) — both rules disclosed, neither test-based. The per-round test trajectories carry the real dynamics: FedAvg’s test ROC-AUC *declines* from 0.974 (round 5) to 0.963 (round 50) — every round except a $+9 \times 10^{-6}$ uptick at the final step — while its validation F1 still improves — client drift unmitigated, with the validation prevalence (2.05%) optimised against the test prevalence (1.31%), so the shipped rule selects the weakest test-ROC monitored round of its own trajectory; FedProx climbs monotonically from 0.931 to 0.975, the proximal term’s stabilisation becoming visible precisely once BatchNorm is removed — a dynamics signature the BN arms’ confound had hidden.

Three boundaries keep the claim honest. (a) PR-AUC, not ROC-AUC, is where work remains: no-BN FedProx’s 0.345 sits below the BN-carrying FedProx-LSTM’s 0.404 and the centralised MLP’s 0.445 — what these arms license is the *attribution*, not a new state-of-the-art row. (b) The paper’s in-partition no-BN control (Section 6.12) is a different substrate — balanced, same-partition — and ends *below* the BN-with-correct-statistics arm there; these raw-substrate arms are its natural-prevalence, out-of-partition complement, and the two results are not readable into each other. (c) FedAvg’s declining test trajectory means its row is a conservative lower bound on what a selection rule with a usable monitor would ship; the sensitivity row and the trajectory are the honest record.

6.5 The BatchNorm-free sequence arms rescue the federated arms

The raw-substrate collapse of Table 6 is a statement about the flattened-feature MLP arms, and the second half of re-run item 8 asks exactly that question: does the same frozen protocol survive natural prevalence when the federated encoder is sequential? The LSTM arms were

executed on owner-side GPU hardware (RTX 4060 Laptop, cuDNN; 2.6 h total wall-clock for both federated arms plus the pooled comparator, against a CPU-only sandbox run that would have been impractically slow—the one experiment where a GPU was a practical requirement rather than a convenience). The substrate is identical by construction: the runner rebuilds the 3D series from the same cached stage-A/B arrays, asserts label equality with the frozen 2D split at startup, runs the 144-statistic feature cross-check whenever the 2D cache is present ($\max|\text{diff}|$ bounded by 5×10^{-3} at float16 storage; on the owner-GPU box the label-equality assertion alone ran, the 2D cache being absent there), and reuses the identical seed-42 Dirichlet shards, so the only degrees of freedom are the encoder and the device. The SolarLSTM (219,265 parameters, consuming the raw 60×24 series directly) is trained under the same Fed-Focal loss, prior-shift calibration, and validation-frozen threshold machinery as every other arm. The one arm still unexecuted in the second owner-GPU batch—SCAFFOLD-LSTM, whose MLP instance is the hardest failure on this substrate—was subsequently run in a second owner-GPU batch together with the seed-43 replication and the SMOTE ablation reported below, under the client-side fixed-step SGD that the control-variate arithmetic requires (the AdamW invalidation of Section 7 is avoided by construction in the federated implementation).

Table 10 reports the outcome alongside each arm’s frozen MLP counterpart, and the asymmetry of Section 6.2 inverts. Centrally, the LSTM is the *weaker* architecture: 0.962/0.314 against the central MLP’s 0.971/0.445 on the same windows. Federated, the ordering reverses. FedAvg-LSTM holds 0.958/0.305 where the FedAvg MLP fell to 0.875/0.056—a +0.083 ROC-AUC and +0.249 PR-AUC swing from the architecture, under identical shards, budgets, and aggregation (the no-BatchNorm control of Section 6.4 attributes the bulk of that swing to BatchNorm-freeness rather than to sequential encoding as such: with the encoder held fixed, the no-BN FedAvg MLP swings +0.088/+0.310). FedProx-LSTM reaches 0.970/0.404: above its own pooled centralised comparator (+0.008 ROC, +0.090 PR) and above XGBoost’s raw PR-AUC (0.329). SCAFFOLD-LSTM lands at 0.970/0.294 against its MLP counterpart’s 0.768/0.057—the largest single architecture swing on this substrate (+0.202 ROC-AUC, +0.238 PR-AUC): the hardest failure of the BatchNorm-carrying matrix is fully recovered by the BatchNorm-free sequence encoder under the identical protocol, completing the four-arm federation matrix on this benchmark. The two contrasts isolate what happened. Between federated arms of identical budget and aggregation, the BatchNorm-free encoder removes the natural-prevalence collapse the BatchNorm MLP suffers — a property the no-BN MLP control shares, so the load-bearing attribute is the architecture’s normalisation-free construction, not the sequence itself; the load-bearing comparison is budget-matched by construction. Between the pooled and federated instances of the same sequence encoder, federation costs nothing (FedAvg) or appears to help (FedProx)—but the latter margin carries a budget caveat, since the central comparator is capped at 30 epochs while the federated arms receive 500 shard passes, and the cleaned-fold parity experiment moved the central MLP by +0.029 ROC-AUC under exactly that correction (Section 5.5).

What does not change is at least as important. No federated arm exceeds the best centralised baseline on either rank metric (FedProx-LSTM sits 0.008/0.044 below logistic regression). The frozen- F_β thresholds transfer badly for three of the four arms: validation F_2 values of 0.727/0.765/0.593 degrade to test values of 0.405/0.309/0.499 under the $1.6\times$ validation-to-test prevalence shift, while SCAFFOLD-LSTM is the exception—its validation F_2 of 0.530 lands at 0.528 on test at a recall-heavy operating point (test recall 0.737)—and the frozen-FPR operating points still realise their targets within band, reproducing (including the same conservative drift at the 0.5% target) the behaviour already documented for the MLP arms. Calibration remains the deployability boundary: all three federated sequence arms beat the raw-substrate climatology floor of 0.013 (FedAvg 0.0112, FedProx 0.0106, SCAFFOLD 0.0120, against every federated MLP arm at or above it, 0.013–0.137), but FedProx-LSTM’s ECE of 0.020 is four times the central LSTM’s 0.005, SCAFFOLD-LSTM carries the worst sequence-encoder calibration (0.035), and the centralised MLP’s Brier of 0.010 remains the best of all ten arms executed on

Table 10: Raw-substrate LSTM arms (frozen protocol, identical substrate, shards, and evaluation as Table 6; owner-GPU execution, seed 42) with each arm’s frozen MLP counterpart repeated from Table 6 for the architecture-versus-federation contrast. Brier and ECE on the calibrated test probabilities at natural 1.31% prevalence. The seed-43 replication and per-client SMOTE ablation of the same arms are reported in Table 11.

Model	LSTM arm		MLP counterpart		Brier	ECE	R@2%FPR
	ROC	PR	ROC	PR			
Centralised	0.962	0.314	0.971	0.445	0.0111	0.005	0.481
FedAvg	0.958	0.305	0.875	0.056	0.0112	0.007	0.431
FedProx	0.970	0.404	0.930	0.149	0.0106	0.020	0.525
SCAFFOLD	0.970	0.294	0.768	0.057	0.0120	0.035	0.521

this substrate. Finally, the FedProx-over-FedAvg stabilisation now appears on every substrate–encoder combination executed (in-partition MLP 0.954 vs 0.575; cleaned fold 0.976 vs 0.906; raw MLP 0.930 vs 0.875; raw LSTM 0.970 vs 0.958, replicated at seed 43 as 0.965 vs 0.961)—the most replicated single finding of this study. The event-level pass (bottom block of Table 7) was executed offline, on CPU, from the GPU run’s stored per-window test probabilities over the same audit metadata—no retraining, 65/65 GOES peaks matched. Where plain FedAvg-MLP detects *zero* of the 65 events — 38 (58.5%) under the diagnostic’s batch-statistics treatment, the only evaluation with a usable validation operating point (Section 6.3) — FedAvg-LSTM detects 35 (53.8%) at 0.2 false-alarm windows per day: the BatchNorm-free rescue survives the validation-frozen thresholds and the cooldown machinery, not just the ranking metrics. FedProx-LSTM is the standout: 49/65 (75.4%) detection at 0.4 false-alarm windows per day, exceeding FedProx-MLP’s raw detection while cutting its false-alarm burden $3.5\times$ (1.3), with a 23.2 h median first-alert lead to flare peak—the joint-longest on this substrate, and (like every arm’s lead at this label geometry) a label-boundary quantity rather than a usable warning time (Table 5). SCAFFOLD-LSTM—measuring, like every SCAFFOLD cell, the sign-inverted, cold-start variant of the five-departure list in Section 4, not the reference algorithm—evaluated in-runner by the same machinery, supplies the sharpest event-level datum of the batch: 55/65 (84.6%) detection at 0.41 false-alarm windows per day and a 23.2 h median lead—its MLP twin’s 83.1% detection rate at one-fifteenth its false-alarm burden, and nominally above the centralised detection frontier (logistic regression 78.5%, centralised MLP 80.0%, both at 0.2–0.3 per day). The crossing is an operating-point purchase rather than a ranking win—it is bought with $1.5\text{--}2\times$ the centralised arms’ alarm budget, the window-level ranking still sits below logistic regression by 0.008/0.154, and the arm carries the worst sequence-encoder calibration (ECE 0.035)—so the data-locality conclusion of Section 6.2 survives at event level with this one priced exception.

6.6 Seed replication and the natural-prevalence SMOTE ablation

The sequence arms executed at a single seed on a single device, and the replication caveat that followed them is discharged by the second owner-GPU batch: all four arms re-executed at seed 43 (model initialisation and the Dirichlet shard draw reseeded, the frozen validation carve kept identical). The ranking structure replicates (Table 11, middle columns)—the seed-42 FedProx/SCAFFOLD statistical tie (0.970 each) resolving into SCAFFOLD’s narrow lead at seed 43. ROC-AUC moves by at most 0.005 on any arm (FedAvg 0.958→0.961, FedProx 0.970→0.965, SCAFFOLD 0.970→0.968, central 0.962→0.962)—inside the device and kernel jitter band predicted in Section 5.6—and every qualitative finding survives: no arm collapses, FedProx stays above FedAvg on both rank metrics, SCAFFOLD stays competitive, and no federated arm approaches the centralised ranking frontier. PR-AUC, as expected at 1.3% prevalence, moves more (FedAvg 0.305→0.271, central 0.314→0.357, SCAFFOLD 0.294→0.368, FedProx 0.404→0.381)

Table 11: Seed-43 replication and per-client SMOTE ablation of the raw-substrate sequence arms (window-level test metrics; protocol, substrate, and evaluation identical to Table 10). Seed 43 reseeds model initialisation and the Dirichlet shard draw; the frozen validation curve is unchanged. SMOTE oversamples each client’s minority class to a 0.25 minority-to-majority ratio (seed 42, flattening the 60×24 series and reshaping).

Model	Seed 42		Seed 43		Seed 42 + SMOTE	
	ROC	PR	ROC	PR	ROC	PR
Centralised	0.962	0.314	0.962	0.357	—	—
FedAvg	0.958	0.305	0.961	0.271	0.961	0.307
FedProx	0.970	0.404	0.965	0.381	0.946	0.288
SCAFFOLD	0.970	0.294	0.968	0.368	—	—

without overturning any load-bearing ordering. Event-level detection moves within a seven-point band at comparable false-alarm rates (FedAvg 35→36, central 41→43, FedProx 49→45, SCAFFOLD 55→50 of 65). The BatchNorm-free sequence-arm rescue is not an artefact of one seed.

The complementary ablation asks whether class rebalancing can substitute for the encoder. Per-client SMOTE, oversampling each shard’s minority to a 0.25 minority-to-majority ratio on the flattened 60×24 series, was run for the two production federated arms at seed 42 (Table 11, right columns). FedAvg is unchanged at window level (0.961/0.307 against 0.958/0.305 natural, with event detection even edging up, 35→40 of 65); FedProx is actively degraded: 0.946/0.288 against 0.970/0.404, a -0.116 PR-AUC loss, with event detection falling 49→37 of 65. The result completes the cleaned-fold SMOTE null (Section 6.15: inert there because the substrate was pre-balanced) at natural prevalence, where it has work to do: oversampling is no substitute for the encoder and costs the strongest arm 0.116 of its 0.404 PR-AUC. This is consistent with the loss-stacking analysis of Section 7—SMOTE, the focal α , and class-rate-aware aggregation all adjust the same effective positive-class weight, and stacking a second balancing mechanism on the protocol’s existing ones overshoots. The rescue is architectural; class balance is neither necessary nor sufficient.

6.7 What survives the leakage-free protocol

With the raw-substrate retraining of Section 6.2 the comparison now spans three substrates—in-partition (overlap disclosed), leakage-free cleaned (rebalanced, synthetic positives), and leakage-free raw (natural prevalence)—and the surviving claims can be tiered across all three. Two claims transfer unchanged to *every* substrate: the *deployability boundary* (only XGBoost’s posteriors survive the prevalence shift on the cleaned fold; on the raw substrate the centralised MLP’s Brier of 0.010 remains the best of all ten arms, with the three federated sequence arms the only federated results to beat the climatology floor (Brier 0.011–0.012 versus 0.013, FedProx-LSTM the best of them at an ECE of 0.020 against the central arms’ 0.005–0.011) and the *threshold-transfer failure* at fixed-FPR operating points. FedProx’s *ROC stabilisation* over plain FedAvg—a composite of genuine early-round optimisation benefit and compatibility with the shipped evaluation path (Sections 6.11–6.12)—transfers to the cleaned fold (0.976 vs 0.906) and, in diminished form, to the raw substrate (0.930 vs 0.875 for the MLP arms; 0.970 vs 0.958 for the sequence arms of Section 6.5; 0.975 vs 0.963 for the no-BatchNorm arms of Section 6.4, where the claim is metric-scoped — the no-BN PR-AUC ordering inverts, FedAvg 0.366 above FedProx 0.345); it is the only claim now observed on every substrate executed, on both encoders wherever both were run, and with and without BatchNorm. Three claims do not transfer beyond the cleaned fold, and all three are statements about the *BatchNorm-carrying MLP* arms: the federated-versus-centralised ROC gap, signed as the federated arm minus its centralised

Table 12: Single-pass test-set evaluation at natural prevalence ($n=331,185$, 1.88% positives), at each model’s F_2 -optimal threshold τ^* selected on *validation*. Calibration: prior-shift (fit on validation). Brier and ECE (15 bins) are computed on the calibrated probabilities. Best value per column in bold; the climatology row is the validation-fitted base rate (the prior-shift convention), and the honest Brier reference is the test-set climatology (0.0185 at 1.88% prevalence), which no arm’s calibrated Brier beats—see Section 6.13.

Model	τ^*	Prec.	Recall	F1	ROC-AUC	PR-AUC	Brier	ECE
Climatology	0.05	0.019	1.000	0.037	0.500	0.019	0.239	0.470
Logistic regression	0.23	0.021	0.967	0.041	0.824	0.235	0.659	0.732
Centralised MLP	0.32	0.030	0.999	0.059	0.888	0.153	0.277	0.439
XGBoost (centralised ref.)	0.37	0.146	0.968	0.254	0.977	0.493	0.063	0.115
FedAvg MLP (federated)	0.36	0.020	0.997	0.038	0.575	0.199	0.136	0.342
FedProx MLP (federated)	0.51	0.021	0.999	0.042	0.954	0.307	0.264	0.496

reference (FedProx +0.003 against the centralised MLP on the cleaned fold; −0.048 against the best centralised arm, logistic regression, on raw; −0.008 for FedProx-LSTM below the same reference; +0.004 past zero against the centralised MLP for the no-BN FedProx), FedProx’s in-partition PR-AUC edge (its raw PR-AUC of 0.149 is a third of the centralised MLP’s 0.445, though the sequence arm reaches 0.404), and the absolute discrimination levels of the flattened-feature federated arms themselves (FedAvg-MLP detects zero of 65 events on the raw substrate — 38 under the diagnostic’s batch-statistics treatment, the only one with a usable validation operating point — while the sequence arm detects 35 and the no-BatchNorm MLP restores the window-level ranking, Sections 6.5 and 6.4). The multi-seed, sweep, ablation, and client-level studies of the following subsections were executed under the in-partition protocol and are retained as protocol-stability evidence; their leakage-free and raw-substrate counterparts remain queued, though the sequence arms’ seed replication is now discharged (Section 6.6; fold-replication remains owner-side compute).

6.8 Headline comparison on the held-out test set (in-partition protocol)

Table 12 contains four findings. **(i) The architecture-controlled comparison now exists.** The v2.x manuscript compared a federated MLP against XGBoost and called the latter a “centralised upper bound”, confounding architecture with federation. With the centralised MLP of the identical 29,377-parameter architecture, the comparison isolates the effect of federation from the effect of architecture—but *not* from optimisation budget (Section 5.5: the federated arm consumed the larger budget, 500 effective local passes versus at most 30 pooled passes). Under that caveat, FedProx (ROC-AUC 0.954, PR-AUC 0.307) did not underperform its pooled counterpart (0.888, 0.153) on this benchmark and configuration—an observation consistent with non-IID federated training plus a proximal regulariser acting as an implicit regulariser for this small MLP, and which we do *not* generalise to the claim that federation is free of cost. **(ii) The shipped evaluation separates FedAvg and FedProx, but the separation is largely an evaluation-path property.** Plain FedAvg’s shipped trajectory drifts (Section 6.10) and its best-validation checkpoint reaches 0.575 ROC-AUC as evaluated; Sections 6.11–6.12 show the same weights carry 0.951 under correct normalisation and that neither algorithm diverges without BatchNorm, so the proximal term is the difference between compatibility and incompatibility with the shipped evaluation path *under the frozen-protocol configuration*, plus a modest early-round optimisation benefit. **(iii) Against the strongest reference,** FedProx retains 97.7% of XGBoost’s ROC-AUC and 62% of its PR-AUC. **(iv) The τ^* column exposes a protocol-honest failure:** thresholds optimised for F_2 at the validation prevalence (48.9%) transfer poorly to the 1.88% test prevalence—all neural models fire almost unconditionally (recall ≈ 1.0 , precision $\approx$ the base rate), and the calibration columns (Brier, ECE) show every




Figure 3: ROC curves on the held-out test set. FedProx (AUC 0.954) tracks the centralised XGBoost reference (0.977) across the low-false-positive region; the centralised MLP (0.888) sits between them; FedAvg (0.575) evaluates far below its weights’ actual information content (Section 6.11: 0.951 under correct normalisation).

neural model *worse calibrated than the climatology*. We analyse this threshold-transfer failure explicitly in Section 6.13 and report prevalence-robust operating points instead.

6.9 Prevalence-robust discrimination: ROC and PR analysis

Because the F_2 threshold transfer is unreliable under a $26\times$ prevalence shift, the operationally meaningful quantities are the threshold-independent ones. At fixed false-positive budgets—the regime a grid operations desk actually inhabits—recall on the test set is:

Three readings matter. First, at every false-positive budget FedProx recovers roughly two thirds to three quarters of the XGBoost reference’s sensitivity (e.g. 0.506 vs 0.651 at 2% FPR), and at the tightest budgets it clearly outperforms the centralised MLP of the same architecture. Second, the PR-AUC column of Table 12 shows the same ordering under prevalence-sensitive scoring (XGBoost 0.493, FedProx 0.307, FedAvg 0.199): in the high-precision region the federated model pays a real ranking price relative to the ensemble, even though it beats its own architecture’s pooled baseline. Third, FedAvg’s recall profile decays fastest with tightening budget—as *evaluated*. Section 6.11 shows this to be a property of the evaluation path, not of the weights: the same checkpoint holds 0.951 ROC-AUC of ranking information under correct normalisation.

6.10 Convergence behaviour: the evaluated drift and its suppression

Figure 4 documents the trajectory that organised this revision’s in-partition narrative for three versions of the paper: plain FedAvg under the six-client Dirichlet partition is not merely slower to converge—as *evaluated by the shipped server-side path*, it drifts from healthy validation AUC

Table 13: Recall at fixed false-positive rates on the test set. These are *discrimination curve statistics*: for each budget, the recall at the point of the test ROC curve where FPR equals the target. They involve no threshold transfer from validation and are therefore robust to monotone miscalibration, but they are *window-level* quantities: the 60-minute sliding windows of one active region are correlated, so a 2% window-level FPR does not translate directly into an operational false-alarm rate per event or per day (event-level evaluation: Section 8).

Model	Rec@0.5%FPR	Rec@1%FPR	Rec@2%FPR	Rec@5%FPR
Logistic regression	0.199	0.287	0.419	0.568
Centralised MLP	0.101	0.147	0.220	0.384
XGBoost	0.346	0.484	0.651	0.864
(centralised ref.)				
FedAvg MLP	0.164	0.247	0.380	0.443
FedProx MLP	0.221	0.345	0.506	0.712




Figure 4: Validation-monitored convergence over 50 rounds: per-round validation F1 (left panel) and validation Recall (right panel) at the development threshold, as evaluated by the shipped server-side path (FedAvg vs FedProx; ROC-AUC values quoted in the text come from the round-level validation history in the run artefacts, not from this figure). Plain FedAvg’s evaluated F1 collapses toward zero after round 20 while its recall saturates at 1.0 (the all-positive degenerate regime); FedProx ($\mu=0.01$) holds its monitored statistics throughout. Section 6.11 shows this trajectory measures normalisation compatibility, not information content: the same FedAvg weights carry ~ 0.95 test ROC-AUC under correct normalisation at every round.

(≈ 0.99) to below-chance ranking (AUC 0.05–0.40) after round 20, while FedProx ($\mu = 0.01$) holds 0.99 ± 0.01 for all fifty rounds. Two protocol details made the observation defensible rather than accidental: monitoring happens on validation only (an earlier revision monitored the test set, which both leaked information and masked the drift), and the best-validation checkpoint rule restored the round-20 FedAvg state. The diagnostic of the next two subsections, however, re-attributes the phenomenon: the drift is a property of the *evaluation path*—BatchNorm statistics that the implementation never transports—and the proximal term’s stabilisation is, in large part, stabilisation of *compatibility with that evaluation path*. The multiseed evidence (Section 6.14) shows the shipped-eval drift is reproducible across seeds; it does not distinguish the two attributions, because every seed was evaluated the same way.

Table 14: Test ROC-AUC / PR-AUC of the identical FedAvg global weights under four BatchNorm treatments (frozen protocol, seed 42; faithful re-run, `bn_diagnostic.json`). Round 20 is the state the shipped checkpoint rule returned. Arm D is what a buffer-transporting implementation would have evaluated.

BN treatment	Round 5	Round 20	Round 50
A: init buffers (shipped)	0.950 / 0.264	0.575 / 0.199	0.098 / 0.010
B: recalibrated (pooled)	0.786 / 0.044	0.893 / 0.099	0.881 / 0.086
C: batch statistics	0.946 / 0.328	0.951 / 0.329	0.953 / 0.288
D: transported round buffers	0.885 / 0.157	0.931 / 0.194	0.935 / 0.209

6.11 BatchNorm-transport diagnostic: decomposing the collapse

The client models contain three `BatchNorm1d` layers, and local training updates their running statistics with every batch. The federation, however, transports *parameters only*: `get_weights/set_weights` move `model.parameters()`, and running statistics are buffers, not parameters. The server-side global model therefore evaluates with whatever buffers it holds—and since the server never executes training-mode forward passes, those buffers remain at initialisation forever. Both shipped checkpoints confirm this directly: every BN layer carries `num_batches_tracked = 0`, `running_mean = 0`, `running_var = 1` after the full 50-round run. In eval mode with init buffers, each BN layer degenerates to the fixed affine map $x\gamma + \beta$; the further the trained weights drift from the region where pre-BN activations are standardised, the less valid that map becomes. This is not FedBN-style per-client local statistics—nothing is averaged, nothing is transported; the statistics are simply absent.

The saturated validation monitor is the second half of the mechanism. With init buffers, the evaluated probabilities squeeze into a narrow band around the 0.35 development threshold (round 20 on validation: $[0.337, 0.374]$; round 50: $[0.219, 0.232]$), so predictions are all-positive (then all-negative) and validation F_1 pins to the closed-form all-positive value $2p/(1+p) = 0.6565$ at the validation prevalence $p = 0.4887$. Because the strict-inequality best-checkpoint rule restores the first round that exceeds the tie, and precision drifted marginally upward (0.489 to 0.491) while the ranking collapsed, the rule restored the *round-20* state—where validation ROC-AUC had already fallen from 0.993 to 0.541—over the healthy round-5 state. FedProx’s monitor tied at 0.6565 from round 5 onward and its ranking never collapsed, so the same rule was harmless for it. The monitor was not merely uninformative; it was adversarial.

Table 14 decomposes the published FedAvg numbers. We re-ran the frozen 50-round FedAvg protocol faithfully (validation trajectory matches the frozen history within ± 0.01 ROC-AUC at every logged round; the checkpoint rule restores round 20 exactly) and evaluated the identical global weights at rounds 5, 20, and 50 under four BN treatments: (A) the shipped init buffers; (B) exact per-client statistics pooled by the law of total variance (an optimistic recalibration oracle); (C) per-batch statistics at evaluation time (adaptive normalisation; a diagnostic of information content, not a deployable protocol); and (D) the element-wise weighted average of the running statistics that local training actually produced in that round—precisely what a buffer-transporting implementation would evaluate.

Three facts follow. First, *FedAvg’s weights never lost discriminative content*: under batch statistics the round-50 state—whose shipped evaluation is 0.098, below chance—holds 0.953 ROC-AUC and 0.288 PR-AUC, and arm D holds 0.935 / 0.209. The published 0.575 is a statement about the round-20 weights’ compatibility with init buffers, not about what FedAvg learned. Second, *the FedProx artefact is equally protocol-dependent*: its returned round-5 weights score 0.954 / 0.307 under init buffers but 0.757 / 0.040 under pooled recalibration and 0.951 / 0.333 under batch statistics. Under batch statistics the two algorithms’ artefacts are statistically indistinguishable (0.951 vs. 0.951 ROC-AUC; 0.329 vs. 0.333 PR-AUC). Third, arm B degrades

Table 15: No-BatchNorm control (frozen protocol, seed 42, 50 rounds, BatchNorm→Identity): test ROC-AUC / PR-AUC at selected rounds plus the shipped best-validation checkpoint. Neither algorithm diverges; the shipped FedAvg–FedProx gap does not survive BN removal.

Algorithm	Round 5	Round 20	Round 50	Selected ckpt
FedAvg, no BN	0.913 / 0.199	0.874 / 0.148	0.873 / 0.108	0.858 / 0.108 (r40)
FedProx, no BN	0.949 / 0.286	0.900 / 0.175	0.871 / 0.128	0.871 / 0.128 (r50)

both artefacts relative to arms A and C—pooled statistics inflate variance across heterogeneous clients, so no single pooled normalisation matches any deployment distribution—which is why arm D (the transport an implementation would actually perform) is the relevant fix and recovers 0.931–0.935 ROC-AUC for FedAvg.

The reading we freeze is that three independent mechanisms—untransported BN statistics, a thresholded metric that saturates under them, and a checkpoint rule that consumes the saturated metric—composed to manufacture a large, reproducible, spurious algorithm gap. Each is individually mundane; their composition produced the single most consequential “finding” of this paper’s in-partition narrative for three revisions. We report this as an evaluation-protocol failure mode in its own right: FL evaluation stacks that transport parameters but not buffers, monitor with thresholded metrics, and select checkpoints on those metrics can manufacture algorithm comparisons that survive seed replication and ablation grids.

6.12 Controlled no-BatchNorm experiment

The diagnostic above establishes that the *evaluated* collapse is an artefact, but not whether the *training dynamics* themselves depend on the BN implementation. We therefore re-ran the identical frozen protocol (data, partition, seeds, rounds, loss, aggregation, monitoring, and checkpoint rule) with the three `BatchNorm1d` layers replaced by identity (28,929 parameters vs. 29,377; `nobn_control.json`). With no BN layers there are no buffers, weight transport is exact, and the shipped evaluation path is the ground truth. Per-round test evaluation is disclosed here as a controlled diagnostic; no test-based selection occurs—the shipped best-validation rule alone determines the reported artefact.

The answer is unambiguous: *without BN, neither algorithm diverges*. Validation F_1 sits at 0.97–0.98 for both algorithms at every logged round (the all-positive tie disappears; the monitor becomes informative), validation ROC-AUC stays at 0.999, and the test trajectories decay mildly (FedAvg 0.913 → 0.873; FedProx 0.949 → 0.871) rather than collapsing. The shipped 0.379 ROC-AUC gap between the two artefacts shrinks to at most 0.036 at any round and 0.002 at round 50. FedProx retains a genuine but modest early-round benefit (round 5: 0.949 vs. 0.913 ROC-AUC, 0.286 vs. 0.199 PR-AUC), consistent with the proximal term improving early optimisation; that benefit then decays at the same rate as FedAvg’s. Meanwhile the BN architecture itself is genuinely useful under correct normalisation (arm C: 0.329 PR-AUC at round 20 vs. 0.148–0.175 no-BN at round 20), so the correct engineering conclusion is BN plus buffer transport, not BN removal. Mild test-side decay in both no-BN arms confirms that client drift is real—it degrades generalisation gradually—but it does not, on its own, produce catastrophic divergence under this partition.

6.13 Calibration and the threshold-transfer failure

The validation prevalence (48.9%) and the test prevalence (1.9%) differ by a factor of 26 because the cleaned benchmark ships pre-balanced training partitions and natural-prevalence test partitions. We treat this *prior shift* as a first-class methodological problem of the benchmark, not a nuisance: it is the central obstacle between the current results and deployment-grade

Table 16: Calibration-arm comparison under the $26\times$ prior shift (review context R6). Every arm is fit on the balanced validation set (48.9% prevalence) and applied once to the test set (1.88%). *Val. Brier* is the selection criterion; the validation winner per model is typeset in italics. Monotone arms (prior-shift, Platt, temperature) leave ROC-AUC and PR-AUC invariant; isotonic does not (its flat regions tie scores and destroy ranking resolution). Bold marks the best *test* Brier per model.

Arm	FedProx MLP			Centralised MLP			XGBoost		
	val. Brier	test Brier	test ECE	val. Brier	test Brier	test ECE	val. Brier	test Brier	test ECE
Raw (none)	0.234	0.264	0.496	0.016	0.277	0.440	0.004	0.063	0.116
Prior-shift	0.234	0.276	0.508	0.016	0.284	0.447	0.004	0.065	0.118
Platt	0.026	0.541	0.642	0.012	0.352	0.442	0.004	0.053	0.080
Isotonic	<i>0.025</i>	0.541	0.639	<i>0.012</i>	0.357	0.449	<i>0.003</i>	0.057	0.075
Temperature	0.064	0.784	0.865	0.012	0.334	0.425	0.004	0.064	0.099

probabilities. The frozen protocol forbids using test labels—including test prevalence—for any selection, so the prior-shift calibrator is fit with the *validation* prevalence as its deployment target. The consequence is visible throughout Table 12: the validation-optimal thresholds sit far below the calibrated test probabilities of most samples, so every neural model operates in an almost-always-alert regime, and *the calibration columns make the failure quantitative*: FedProx’s calibrated Brier (0.264) and ECE (0.496) are *worse than the climatology’s* (0.239, 0.470), i.e. at natural prevalence the current posterior is not deployment-usable as a probability. The XGBoost reference suffers the least (Brier 0.063, ECE 0.115), which is why it converts ranking into a meaningful F_1 of 0.254 while the neural models do not.

Three readings follow. First, *rebalanced training data invalidates naive threshold transfer*: an F_2 -optimal threshold at 49% prevalence is semantically meaningless at 1.9%. Second, the correct deployment instrument, pending full recalibration, is the recall-at-fixed-FPR operating points of Table 13, which are invariant to monotone miscalibration and to prevalence. Third, the honest remedies for calibration itself are prior correction toward an *estimated* deployment prevalence (e.g. from unlabelled deployment data, which leaks no labels), Platt/isotonic fitting on a natural-prevalence validation carve-out, or the negative-quantile FPR thresholds of Section 4.8; using the test set’s own statistics, as the v2.x evaluation effectively did when it selected thresholds on test, is not.

Calibration-arm comparison (executed). The released calibration-comparison harness (`run_calibration_comparison.py` in the `experiments` directory) was executed against the trained models of the frozen-protocol run: five arms (raw, prior-shift, Platt, isotonic, temperature), each fit on the balanced validation set, applied exactly once to the test set, with the arm selected by minimum *validation* Brier (a proper scoring rule, so the selection itself is honest). Table 16 reports the outcome. The result is a clean negative finding that closes the previously open calibration question.

Four facts matter. (i) **The selection criterion itself is broken by the prior shift.** Minimum validation Brier selects isotonic for all three models; for both neural models isotonic is the *worst* arm on test (FedProx: 0.541 vs 0.264 raw, i.e. selecting on validation doubles the test error of the posterior). (ii) **No arm rescues the neural posteriors.** Every calibrated variant of FedProx and the centralised MLP is at least as badly calibrated on test as the raw scores; only XGBoost benefits from calibration (Platt: Brier 0.063 $\rightarrow$ 0.053, ECE 0.116 $\rightarrow$ 0.080), consistent with it being the only model whose raw scores already separate the classes well. (iii) **Isotonic damages ranking, not just probabilities.** Its step function ties large score ranges, collapsing PR-AUC from 0.307 to 0.182 (FedProx), 0.153 to 0.099 (MLP), and 0.493 to 0.426 (XGBoost); ROC-AUC and PR-AUC are exactly invariant under the three monotone smooth




Figure 5: Five-seed statistical validation (seeds 42–46, 20 rounds, frozen protocol). (a) Test metrics with 95% confidence intervals. (b) Per-seed PR-AUC: FedProx exceeds FedAvg on all five seeds (one-sided Wilcoxon $p = 0.031$).

arms, as they must be. **(iv) Validation-frozen FPR thresholds fail jointly with probability calibration.** Applying the negative-quantile thresholds of Section 4.8 (frozen on validation at 0.5/1/2/5% targets) to the test set realises 28.2/37.0/50.4/72.1% actual FPR for FedProx (25.5/40.7/51.6/63.6% for the centralised MLP, 20.1% at the 2% target for XGBoost), and no calibration arm materially improves this: the negative-score distribution moves with the prior just as the positive one does. The operational conclusion is unchanged but now empirical rather than hypothetical: at this benchmark’s balanced training prevalence, *no* validation-fit decision layer survives the shift, and deployment-grade probabilities require either natural-prevalence validation data (available only in the raw benchmark) or unlabelled prior estimation at the deployment site.

6.14 Multi-seed statistical validation

The v2.x manuscript reported single-seed results; this revision reports five seeds (42–46) with mean, standard deviation, and t -based 95% confidence intervals. On the test set, FedProx achieves PR-AUC 0.293 ± 0.029 (CI95 ± 0.036) and ROC-AUC 0.946 ± 0.013 (CI95 ± 0.016), against FedAvg’s 0.266 ± 0.025 (CI95 ± 0.031) and 0.916 ± 0.036 (CI95 ± 0.044). FedProx is better than FedAvg on PR-AUC in *five of five seeds* (one-sided Wilcoxon signed-rank $p = 0.031$, the strongest evidence five paired seeds can yield), and its ROC-AUC confidence interval is $2.7\times$ narrower. The F_1 difference is not significant ($p = 0.5$): the threshold-transfer regime dominates F_1 for both algorithms, which is consistent with Section 6.13. What the five-seed study supports is a reproducibility claim about the shipped pipeline: FedProx’s superiority over FedAvg under the frozen-protocol configuration is sign-consistent across every seed and statistically significant. What it does not support is broad generalisation, and—after Sections 6.11 and 6.12—it must be read as a statement about the composite (training $\times$ stale-buffer evaluation $\times$ saturated-monitor checkpointing), not about the algorithms in isolation: every seed was evaluated through the same never-transported BN statistics, and the no-BN control shows the underlying optimisation gap is small. Five seeds is a small sample, the confidence intervals quantify *seed-to-seed variability* of this fixed configuration, not the uncertainty of future observations or of other partitions, and Section 6.15 shows the shipped-eval gap narrows when FedAvg is assisted by distribution-aware aggregation. The appropriate summary is that FedProx was the most consistent stabiliser of the shipped evaluation path among the configurations studied, not that it is load-bearing in every federated flare-prediction deployment.




Figure 6: Component-isolation grid (2 algorithms $\times$ 3 losses $\times$ 2 aggregations $\times$ 2 balancing arms; 20 rounds, frozen protocol, single-pass test). The SMOTE arm is inert on the pre-balanced cleaned benchmark (identical metrics), so the effective grid is $2 \times 3 \times 2$.

6.15 Component attribution: the 24-cell ablation grid

Because the v2.x system trained SMOTE, Fed-Focal loss, distribution-aware aggregation (DA-FL), and FedProx *simultaneously*, a FedProx win could not be attributed to any single component. The ablation grid isolates them (Fig. 6). Four patterns emerge. **(i) FedProx is the most consistent stabiliser, but not the only remedy:** every FedProx cell lands in a narrow ROC-AUC band (0.954–0.959) whereas the FedAvg cells range from 0.880 (focal loss, plain aggregation—the collapsing configuration) up to 0.963 (plain BCE with distribution-aware aggregation). That best FedAvg cell *exceeds* the FedProx headline (0.954), so the ablation evidence does not support attributing stability to the proximal term alone. **(ii) DA-FL partially rescues FedAvg:** with focal loss, switching plain to DA-FL aggregation lifts FedAvg from 0.880 to 0.950, and the best cell overall (0.963) is plain BCE with DA-FL, suggesting distribution-aware weighting mitigates the same client-drift the proximal term suppresses—at the cost of rate metadata sharing. **(iii) Loss choice is secondary at this prevalence:** across cells, focal, weighted-BCE, and plain-BCE arms differ by at most 0.05 PR-AUC, with weighted BCE marginally best for FedAvg-plain (0.308 vs 0.280). **(iv) SMOTE is inert:** the cleaned benchmark’s training partitions are pre-balanced (48.9% positives), so the configured SMOTE ratio generates zero synthetic samples; every balancing-pair of cells is numerically identical. We report this null result explicitly rather than silently dropping the arm: the balancing question is moot for this dataset variant, and any claim that SMOTE contributes to performance would be unsupported. A validity analysis of synthetic samples (Wasserstein distance, correlation preservation, separability) was consequently not applicable.

6.16 Client-level heterogeneity: who benefits from federation?

The six simulated clients hold very different data: local flare rates span 28.6% (Client B) to 80.5% (Client A), and shard sizes span 1,503 (Client E) to 23,418 (Client B). Figure 7 quantifies what each client gets out of the federation under the untouched-holdout protocol: a dedicated re-run in which every client’s shard was split 80/20 *before* any training, FedAvg and FedProx were retrained from scratch on the 80% federation-train portions (same frozen configuration: 50 rounds, six clients, seed 42), per-client local models were trained on the same portions with a matched local budget, and both were evaluated once on the holdouts no model had ever seen. Three findings hold. **(i) FedProx matches local-only training and helps the smallest custodians.** Mean PR-AUC is 0.988 (FedProx global) versus 0.983 (local-only); on the four




Figure 7: Per-client PR-AUC of local-only models versus the federated global models, evaluated on each client’s *untouched* 20% holdout (R14 protocol: holdouts carved before training; both global models retrained on the 80% federation-train portions only, so local and global scores are clean generalisation estimates; neutral client labels, no real institution participates).

largest clients the two are within 0.001–0.010 of each other, while the two smallest shards gain clearly (Client E: 0.965 $\rightarrow$ 0.993, +0.027; Client F: 0.961 $\rightarrow$ 0.980, +0.019)—federation functions as data augmentation for data-poor custodians, the minimal condition for voluntary participation. **(ii) The diverged FedAvg average is worse for everyone.** Its global model scores 0.747–0.973 per client (mean 0.866), below local-only training on *every* client: the collapsed average is worse than staying isolated, exactly as the convergence analysis of Section 6.10 predicts. **(iii) The clean protocol confirms the earlier directional result.** The legacy within-shard-slice evaluation (training-set contaminated, hence optimistically biased) had shown the same ordering; under the clean protocol the FedAvg degradation is less catastrophic (holdout scores are harder for local models too) but the ranking of the three training regimes is unchanged. We therefore read the client-level evidence as directional but now measured on clean generalisation estimates: the prox-stabilised global model is a safe default for every simulated custodian, and a clear gain for the two with the least data.

6.17 Interpretability across models

The v2.x interpretability story relied on the centralised XGBoost surrogate alone; the improved pipeline additionally attributes the *federated* FedProx model (DeepSHAP with a 100-sample background) and cross-checks the two. At the feature level the models agree only moderately (Spearman $\rho = 0.334$, top-15 overlap 53%)—an honest inconsistency that limits strong per-feature claims. At the level of *physical parameter groups*, however, the two attribution structures are consistent: current-helicity parameters (TOTUSJH/ABSNJZH/MEANJZH/TOTUSJZ/MEANJZD) are the largest group for both (39.9% of XGBoost attribution mass, 29.0% for FedProx), followed by the R-value flux-emergence proxy (24.7% / 23.2%) and the Lorentz force (22.7% / 18.8%). Both models therefore concentrate their decisions on the established physics of flare productivity—magnetic free energy, current systems, and non-potential stress [5, 6, 20]—and the federated model is not an uninterpretable artefact of aggregation. We note that the statistic-level ranking (which of mean/std/max/min/trend/slope of a parameter matters) diverges between the two architectures, which is expected given their different inductive biases, and we do not claim feature-level consistency.




Figure 8: Mean absolute SHAP values, top 15 features, for the centralised XGBoost reference on the test set. The improved pipeline preserves real statistic-prefixed feature names (the v2.x pipeline reported generic indices, which this revision fixes).

6.18 Communication cost

The global model has 29,377 parameters (0.117 MB at 32-bit precision). Full-parameter exchange costs 117.5 KB per direction per client per round: 0.235 MB round-trip, 1.41 MB per round across the six-client cohort, and 70.5 MB aggregate over the 50-round run (11.75 MB per client). These are *parameter-exchange arithmetic* from the frozen configuration—bytes counted, not measured on a live deployment—and they quantify the *training cadence* (one round-trip per round) separately from the *observation cadence* (magnetogram records arrive every 12 minutes and the pooled windows are spaced roughly 60 minutes apart; the global model is refreshed per training cycle, while alerts can be emitted per window). Wall-clock communication latency, aggregation compute, client training time, straggler recovery, and secure-aggregation bandwidth overhead were *not* measured: the experiments ran in-process on a single CPU container, with no real network. Enabling secure aggregation would add a bounded multiplicative factor (mask material per client per round), quantified in the deployment notes but not evaluated here (Section 4.9).

Honest scaling against the data actually centralised. The relevant comparison is not raw magnetograms (our pipeline never ships unprocessed data) but the feature arrays the centralised alternative would need to receive once. On the cleaned substrate the entire training feature set is $82,121 \times 144$ float32 values (47.3 MB): the cohort’s 50-round parameter exchange (70.5 MB) is *1.5 times larger*, so no bandwidth saving accrues for a single training cycle at this model size — communication frugality here is a property of the continuous data stream (new records every 12 minutes) and repeated retraining cycles, not of one run. On the raw 3D substrate the ordering reverses: centralising the training windows ($214,888 \times 60 \times 24$ float32, ≈ 1.24 GB) exceeds the LSTM cohort’s lifetime parameter exchange at 32-bit round-trip precision (219,265 parameters, ≈ 526 MB) by a factor of 2.4. The often-quoted ≈ 130 MB figure holds only under an fp16 uplink-only convention that the code base lists as future compression work, not as an

implemented path; we state the 32-bit round-trip number as the honest default. Latency and availability engineering remain deployment work.

6.19 Data-locality cost analysis

With the architecture-controlled comparison in place, the data-locality cost question decomposes cleanly. Relative to the *same-architecture* centralised MLP, the federated model did not underperform and in fact scored higher (ROC-AUC 0.954 vs 0.888; PR-AUC 0.307 vs 0.153) on this benchmark and configuration—an observation we report with its two caveats: the federated arm consumed the larger optimisation budget (Section 5.5), and the mechanism is consistent with a regularisation-like effect rather than a benefit of federation as such. At this benchmark’s scale, six disjoint shards with 29k parameters and a proximal regulariser behave like an ensemble-like regulariser. Relative to the *strongest available* centralised reference (XGBoost, 300 trees), FedProx retains 97.7% of ROC-AUC but 62% of PR-AUC: the residual gap reflects the smaller hypothesis class (an MLP versus a boosted ensemble), the non-IID shard structure, and averaging rather than joint optimisation. The honest summary for a deployment decision: federation under this protocol cost no ranking skill relative to a pooled MLP of the same class *in this study*; no raw magnetogram left any simulated client; and the remaining gap to the tree-ensemble reference is an architecture-budget question, not a data-locality question. What the study does *not* establish is that these in-partition numbers generalise: budget parity and natural-prevalence retraining have both since been executed (the gap survives budget parity, Section 5.5, but is substrate- and BatchNorm-conditional, Sections 6.2–6.4), and within-fold region-disjoint validation remains the open item (Section 8).

7 Discussion

7.1 What the provenance audit changes, and what it leaves intact

The data-provenance audit of Section 5.7 is, in retrospect, the single most consequential methodological step of this study: the in-partition numbers that dominate the literature’s usage of the cleaned benchmark—including this study’s own headline until this revision—evaluate on windows that the training pool already contains, most completely for the minority class. Two lessons follow. First, *pipeline-level leakage auditing is not enough*: the runtime audit of this framework verified disjoint shards, immutable splits, and single-pass test evaluation, yet could not see that the dataset’s own train/test exports share instances; only raw-metadata alignment exposed it. Benchmark consumers who cannot verify provenance at the instance level should treat same-partition numbers as upper bounds. Second, *the honest comparison is less flattering but more informative*: on the leakage-free fold the federated-versus-centralised gap vanishes, and what separates the arms is no longer discrimination (everyone exceeds 0.906) but calibration and false-alarm burden—dimensions on which the neural arms, federated or not, lose to the gradient-boosted reference. The federated story that survives is the narrower one this section has always argued: under adversarial heterogeneity, proximal regularisation keeps the federation converging where plain averaging collapses (0.976 vs 0.906, with checkpoint-independent dynamics), and it does so without data movement. On the raw substrate the same statement is BatchNorm-conditional: every BatchNorm-carrying MLP arm collapses there under every regulariser studied, while every BatchNorm-free arm — the sequence arms of Section 6.5 (FedProx-LSTM 0.970 vs FedAvg-LSTM 0.958) and the no-BatchNorm MLP control of Section 6.4 (FedProx 0.975 vs FedAvg 0.963) — restores it, relocating the fragility from the federation mechanism to the BatchNorm-under-federation interaction. The stabilisation reading itself gains a sharper, metric-scoped form there: with BatchNorm removed, FedProx’s test ROC-AUC climbs monotonically (0.931 to 0.975 over 50 rounds) while FedAvg’s declines (0.974 to 0.963) despite its improving

validation monitor — the proximal term’s variance reduction is visible exactly once the confound is gone — but on PR-AUC the no-BN ordering inverts (FedAvg 0.366 above FedProx 0.345), so the raw-substrate stabilisation claim is ROC-scoped.

7.2 Why the shipped evaluation separates FedAvg from FedProx

The asymmetry between FedAvg and FedProx in Table 12 deserves a mechanistic reading, and the present revision’s diagnostics give it a different centre of gravity than earlier revisions assumed. Under the realised partition (Fig. 2) clients do receive genuinely different problems (one holds 36% of the data at a 29% flare rate, another 2% at 43%), and client drift is real—the no-BN control shows both algorithms decaying gently in test quality over rounds. But the catastrophic component of the shipped asymmetry is an *evaluation-path* effect: the implementation transports parameters only, so BatchNorm running statistics remain at initialisation on the server, the evaluated FedAvg probabilities squeeze toward the threshold, the all-positive F_1 tie follows, and the checkpoint rule restores a collapse-round state (Sections 6.11–6.12). FedProx’s proximal term (Eq. (4)) anchors each client’s optimisation within a trust region around the broadcast model; a side effect is that the weights stay near the initialisation whose pre-BN activations the (0, 1) buffers normalise correctly, so the shipped evaluation remains valid for it. The in-partition dynamics reproduce leakage-free in direction if not magnitude: on the temporally-preceding fold FedAvg’s ranking survives via its round-35 validation checkpoint (0.906) while its final-round state has collapsed, whereas FedProx’s evaluated ranking never degrades (Section 6.1). The multi-seed study of Section 6.14 shows FedProx above FedAvg on all five seeds (one-sided Wilcoxon $p = 0.031$) with a $2.7\times$ narrower ROC-AUC confidence interval, and the 24-cell ablation grid shows every FedProx cell stable (0.954–0.959) while the FedAvg cells range down to 0.880—all under the same evaluation path, so these reproduce the composite, not the pure algorithm contrast. The defensible statement is narrower than “FedProx is load-bearing”: the proximal term is the *most consistent stabiliser of the shipped evaluation path* among the components studied (distribution-aware aggregation also stabilises the shipped FedAvg cells, best 0.963), it retains a modest genuine optimisation benefit at early rounds, and any deployment of this codebase should transport BN statistics (or remove BN) before drawing algorithm conclusions at all. The raw-substrate LSTM arms extend that evidence across encoders: FedProx-LSTM 0.970 versus FedAvg-LSTM 0.958 on both rank metrics, with identical shards, budgets, and aggregation, makes the stabilisation the only finding of this study observed on every substrate executed and on both encoders wherever both were run. The architecture result itself—the sequence encoder removing the natural-prevalence collapse the MLP suffers—is measured as a contrast, not explained as a mechanism: reading the 60-step temporal profile directly plausibly matters more than any pooled 144-statistic summary when every client’s local positive rate is skewed, but isolating *why* (temporal separability versus effective capacity versus optimisation noise) would require ablations we leave queued.

7.3 Operational positioning: honestly, none

This paper makes no operational claim. The framework does not replace NOAA/SWPC or other national warning centres, does not itself predict GICs, and — on the evidence of its own calibration tables — does not currently emit probabilities that could responsibly be consumed as alerting inputs: no validation-fit calibrator or threshold survives the prevalence shift, and the frozen false-positive budgets are violated by an order of magnitude at test time (Section 6.13). Converting flare probabilities into geomagnetic-hazard and GIC guidance would require the downstream pipeline — CME detection and interplanetary propagation modelling, storm forecasting, and power-system response modelling — none of which is evaluated here; flare class is neither necessary nor sufficient for geoeffectiveness. As a window-level ranking statistic only,

Table 13 describes what the federated models contribute: at a 2% window-level false-positive budget, FedProx’s ranking recovers 51% of positive windows (recall 0.506). Two cautions apply before reading that as an operational detection rate: the figure is a window-level curve statistic over correlated sliding windows, not an event-level detection rate (the event-level evaluations on both substrates are reported in Tables 5 and 7), and the decision threshold required to operate at that FPR budget must be calibrated on natural-prevalence data before deployment—the re-executed frozen-FPR report quantifies how far short validation-frozen thresholds fall, realising 20–50% test FPR against the 2% target across models (Section 6.13). The compute footprint is worth stating as an engineering fact rather than a deployment promise: the entire training programme converges within 50 rounds on a single CPU and fits inside the budget of one workstation, and the benchmark spans a single solar cycle (cycle 24), so any retraining cadence claim would additionally need to establish skill stability across cycle phases, which this single-cycle study cannot. No SCADA, grid-management, or control-room interface is claimed; the decision-threshold boundary in Fig. 1 is drawn to make explicit what is *not* provided (a deployable calibrated alerting layer), not to promise one.

7.4 Engineering lessons for first FL deployments in the physical sciences

The released framework carries a multi-version stability narrative (summarised in Appendix A), and its diagnostic arc is independently useful to any group attempting a first FL deployment on imbalanced scientific data. Five lessons generalise. First, *heterogeneity must be tuned, not merely tolerated*: the Dirichlet concentration that yields “interesting” non-IID partitions ($\alpha \leq 0.5$) can also yield pathological ones (0.3%–99.8% local class rates) under which nothing converges; the stable regime for this benchmark sits near $\alpha = 1.0$. Second, *imbalance handling interacts across layers*: SMOTE, focal α , class-rate-aware aggregation, and threshold selection all adjust the same effective positive-class weight, and uncoordinated adjustments (our focal α was at one point allowed to drift to 0.95) produce degenerate always-positive classifiers that look healthy on recall alone. Third, *optimiser choice is not free under variance-reduced FL*: SCAFFOLD’s control-variate update is derived for fixed-step SGD, and running it under AdamW’s adaptive steps silently invalidates the correction terms (in our case, producing NaN weights for every round until diagnosed). Our own shipped variant carried a quieter direction error of this class — the sign-inverted client control-variate update disclosed in Section 4, which trained plausibly for fifty rounds before a line-by-line audit found it — so the lesson generalises: direction errors in correction terms do not announce themselves, and neither magnitude-level disclosures nor a passing battery substitute for reading the update against its reference equation. Fourth, *feature-name hygiene is part of scientific reproducibility*: an earlier revision’s generic `feature_N` naming in the statistical representation is why its SHAP layer needed an analytical decoding pass (Fig. 8); the improved pipeline preserves statistic-parameter names end-to-end, which is precisely what makes its interpretability auditable. Fifth—the lesson of the present audit, and the one we would put first for any evaluation-heavy project—*audit the evaluation stack itself*: parameters-only weight transport, a thresholded monitor, and strict-inequality checkpointing are each mundane engineering choices, and their composition manufactured a seed-reproducible algorithm gap that survived three revisions of this paper, a five-seed statistical study, and a 24-cell ablation grid. The diagnostic that broke it open cost minutes, not GPU-weeks (Section 6.11): evaluate the same weights under corrected normalisation before believing any instability narrative.

8 Limitations and Future Work

Privacy and security scope. This work enforces the data-locality constraint (raw data never leaves the client), which is the binding constraint of the simulated setting, but the stronger cryp-

tographic notion of privacy is not yet evaluated end-to-end. The communication transcript of Eq. (3)—gradients or weight deltas—is known to leak training information under gradient-inversion attacks [29], particularly for over-parameterised models on low-entropy data. The full threat model (curious server, gradient inversion, membership inference, model poisoning, impersonation, interception) is stated in Table 1: secure aggregation is implemented as a numerical sketch whose tests verify arithmetic round-trip cancellation only (deterministic mask keys; no cryptographic guarantee), differential-privacy hooks are provided, and poisoning/backdoor defence is not addressed at all. The accuracy cost of enabling secure aggregation and differential privacy on this benchmark has not been measured; before any cross-custodian deployment handling genuinely sensitive observations, that budget must be re-measured. Until then, “parameter-transport-only”, not “cryptographically private”, is the correct description of the verified property.

Region-identity, validation inflation, and the resolved provenance question. An earlier revision of this paragraph stated that the cleaned export lacks HARP identifiers and that the definitive split fix awaited the raw benchmark. That dependency is now resolved: every cleaned window has been aligned to its raw instance (Section 5.7), and the audit inverted the concern—the material leakage is not validation-region overlap but instance sharing between the same partitions’ train and test exports, now quantified (100% of flaring test windows; 56,005 of 56,005 unique verified train instances, 56,181 rows) and removed by the temporally-preceding fold executed in this revision. The region-disjoint split generator remains available for within-fold validation studies, and no region spans partitions, so every temporally-preceding fold is automatically region-disjoint. What remains open: the remaining three folds, leakage-free multi-seed and ablation replication, and the synthetic-share question—85.7–90% of training positives are TimeGAN-generated windows, a substrate property that all published results on this export inherit and that raw-benchmark retraining would have to revisit—the raw-substrate contrast of Section 6.2 addresses it only indirectly (its positives are raw by construction).

Threshold transfer under prevalence shift. The frozen protocol selects F_2 -optimal thresholds on validation (48.9% prevalence, mirroring the balanced training partitions), but deployment prevalence is 1.9%: the transferred thresholds put every neural model in an almost-always-alert regime, and the calibration comparison (Section 6.13) shows the neural posteriors are currently *worse than climatology* on Brier and ECE at natural prevalence. We deliberately did *not* repair this by touching test statistics—the v2.x evaluation effectively did, and that is protocol leakage. The remedies are the frozen-FPR operating-point protocol (implemented and unit-tested), prior correction using an *estimated* deployment prevalence from unlabelled data, or a natural-prevalence validation carve-out (the harness for all of which is implemented and smoke-validated). The natural-prevalence re-run has since landed (Section 6.2): its validation carve is likewise natural, so its decision layer does not sit in the almost-always-alert regime, though its frozen F_2 thresholds still transfer imperfectly under the residual $1.6\times$ shift. The in-partition threshold-conditional test metrics (precision, recall, F1) of the balanced substrates should still be read as documenting the decision layer’s failure mode, not as skill estimates, and no probability from the neural models of the balanced-prevalence substrates should be treated as deployment-grade.

Client-level evaluation semantics. The legacy per-client comparison evaluated each client’s local model on a 20% holdout of its own shard while the federated global models trained on the *full* shards, so the global models’ per-client scores were within-distribution rather than clean generalisation estimates, i.e. optimistic for the federated side. That concern is discharged: the untouched-holdout evaluation—every client’s shard split 80/20 *before* training, both federated

arms retrained from scratch on the 80% portions, and all models evaluated once on holdouts no model had seen—is executed and is the protocol behind Figure 7 and Section 6.16. What remains open on this axis is substrate breadth: the executed study covers the cleaned, in-partition setting, and its raw-substrate counterpart is queued.

Partition realism. Shard disjointness is now guaranteed by construction (index-level allocation, audited at runtime), but the partition is still *statistical* label skew, not physical observatory coverage: no mapping from active regions to the agencies whose instruments actually observed them exists in the cleaned export. The partitioning module supports a geographic mode that activates when region metadata is available; until then, the six-client federation remains a simulation of data locality, not a replication of it.

Algorithm coverage. The MLP path is now evaluated on all three substrates, and the LSTM path—which consumes the raw 60×24 windows and matches the temporal architectures that lead centralised SWAN-SF benchmarks [22]—has been evaluated federated on the raw substrate (Section 6.5, owner GPU, two seeds): the sequence arms do not collapse at natural prevalence, at window level and at event level, retroactively validating them as the route to closing the tree-ensemble gap in federated settings. The SCAFFOLD-LSTM arm has since been executed (0.970/0.294 window-level; 55/65 events at 0.41 false-alarm windows per day—both cells measuring the sign-inverted, cold-start implementation case of Section 4, not the reference algorithm), recovering the hardest raw-substrate MLP failure, and a second-seed replication confirms every arm’s ROC-AUC within ± 0.005 (Section 6.6). What remains on this axis: their cleaned-fold pairing. Personalised FL (per-client heads over a shared backbone) is a further direction: agencies may prefer a global model for situational awareness plus a locally adapted model for their own asset fleet.

Event-level evaluation. All headline metrics in this paper are *window-level*: every 60-minute sliding window is scored independently, yet windows overlapping the same active region are strongly correlated, and one physical flare event triggers many correlated decisions. A 2% window-level FPR therefore does not mean 2% of operational alerts are false, and a detected event can generate a burst of duplicate alerts. The event-level evaluation module (detection rate, missed events, alerts per detected event, duplicate-alert rate with cooldown deduplication, warning lead time, alert rate per day) is implemented, unit-tested, and—after the raw-metadata audit recovered the event grouping keys and GOES peak times—executed on the trained models of both substrates (Tables 5 and 7). The window-level recall and FPR figures elsewhere in this paper still carry the window-semantics caveat stated in Table 13’s caption, and the event-level results themselves are single-fold.

Benchmark-construction effects and the raw-benchmark experiment. The cleaned variant’s training partitions are pre-balanced (RUS, Tomek links, TimeGAN augmentation), which (i) renders the SMOTE balancing arm inert—reported as a null result—and (ii) creates the $26\times$ prevalence shift that drives the calibration story. Because the cleaned variant bundles several interventions (imbalance handling, imputation, normalisation, synthetic augmentation), the in-partition results answered the question “can a federation learn the cleaned benchmark?” rather than “can a federation learn the underlying problem?”. This extension has now been *executed* (Section 6.2): the frozen protocol was retrained on the raw, unbalanced partitions with the release’s imputation and normalisation stages reproduced in-pipeline from their published description and every parameter fitted on the training pool only, with the reproduction verified against the released export on the audit-matched windows (median Spearman $\rho = 0.898$ at observed-nonzero positions). The outcome reassigns the federated versus centralised question

entirely: the centralised arms are substrate-robust (logistic regression unchanged to three decimals; centralised MLP PR-AUC 0.382→0.445), XGBoost’s minority-class precision was inflated by the synthetic positives (PR-AUC 0.457→0.329), and every federated *MLP* arm degrades sharply on the natural-prevalence substrate (FedProx 0.976/0.308→0.930/0.149; FedAvg PR-AUC 0.056 and zero events detected; SCAFFOLD, executed here for the first time on a real substrate, fails at 0.768/0.057 — in its disclosed sign-inverted, cold-start implementation, Section 4) —while the sequence arms of Section 6.5 substantially restore the federated standing (FedProx-LSTM 0.970/0.404), relocating the fragility to the BatchNorm-under-federation interaction: the no-BatchNorm control of Section 6.4 removes the three BN layers and nothing else, and the federated MLP recovers to at-or-above the sequence arms’ ranking. A v4.8 errata qualifier on the MLP half of that sentence was discharged at v4.9: the raw-substrate federated MLP evaluations pass through the same untransported-BatchNorm path that Section 6.11 shows manufactures apparent in-partition collapse (parameters transport, running buffers do not), and the checkpoint-level diagnostic has now run over that substrate (Section 6.3) — its verdict is that the degradation is *not* BN-evaluation-conditional (pooled recalibration destroys, batch statistics do not recover). The qualifier that remains is provenance-level, and it is disclosed: the v3.x-era checkpoints were never persisted, so the diagnostic executed over a fresh seed-42 replication whose arm-A numbers deviate from the published rows by 0.002–0.023 ROC-AUC (a recorded MISMATCH, frozen and battery-pinned with the artefact); the sequence arms carry no BatchNorm and are unaffected, and the no-BatchNorm re-training has since run over that substrate and completes the attribution: with the three BN layers removed and nothing else changed, both federated MLP arms recover to at-or-above the sequence arms’ ranking (FedAvg 0.963/0.366, FedProx 0.975/0.345, the latter above the centralised MLP on ROC-AUC, Section 6.4), so the MLP half of this paragraph’s framing is superseded by the BatchNorm-under-federation attribution; the SCAFFOLD arm has no no-BN counterpart in the executed register, so its reading stays qualified. What remains open is fold-replication breadth: the seed-replication and natural-prevalence balancing questions queued in earlier revisions are now executed (Section 6.6: every arm’s ROC within ± 0.005 at seed 43; per-client SMOTE is a clean negative—FedAvg unchanged, FedProx -0.116 PR-AUC—so the rescue is architectural rather than class-balance-driven).

Executed and queued re-run programme. For auditability we collect the re-run experiments in one place. Five items of the original programme have since been *executed* in the independent re-execution (Section 5.6) and are reported in Section 6: the validation-frozen FPR operating-point report (realised test behaviour, quantified in Section 6.13); the calibration-comparison harness across raw / prior-shift / Platt / isotonic / temperature arms (Table 16); the untouched-holdout client evaluation (Section 6.16); the budget-matched centralised retraining at a 500-epoch cap (Section 5.5: the shipped reference was undertrained—early stopping fires at epoch 13—and at full budget parity the centralised MLP reaches 0.917/0.219 versus FedProx’s 0.954/0.307, so the headline gap survives budget parity); and the $\alpha=5.0$ promotion run at the full 50-round protocol (negative: 0.955/0.296 versus 0.954/0.307; $\alpha=1.0$ retained, with the incidental finding that plain FedAvg’s shipped evaluation does not collapse at $\alpha=5.0$ —a severity effect of the evaluation-path drift dissected in Section 6.11). The raw-metadata provenance audit subsequently resolved the metadata blockers and executed items (1)–(3) on the leakage-free fold (Section 6.1: region-disjointness holds by construction of the temporally-preceding fold; natural-prevalence test boundary; event-level evaluation with 65/65 GOES peak matches), and the raw-substrate retraining executed item (5) together with the SCAFFOLD half of item (4) on the natural-prevalence benchmark (Section 6.2), and the LSTM half of item (4) was executed on the owner GPU (Section 6.5: the sequence arms restore the federated standing, FedProx-LSTM 0.970/0.404, at 2.6 h wall-clock on the owner GPU — the CPU-only protocol would have been impractically slow), and the second owner-GPU batch executed the remaining queue (Section 6.6: SCAFFOLD-LSTM, the seed-43 replication of all four arms, and

the natural-prevalence SMOTE ablation, 8.7 h, validated by a 64-check artefact audit), and the third owner-GPU batch executed the Round-13 register’s two decisive controls (the arm-B centralised sanity check, seconds, and the no-BatchNorm raw-substrate federation, 1,699 s; artefacts committed and battery-pinned, Sections 6.3 and 6.4). What remains queued is fold-replication breadth: the remaining temporally-preceding folds for the leakage-free and raw-substrate runs, and the sequence arms’ cleaned-fold pairing. Items previously labelled operational preconditions have been discharged; the event-level figures of Tables 5 and 7 are the operational performance figures of this revision, with the fold-replication caveat stated above.

9 Conclusion

We presented an audit of federated solar-flare prediction on the public SWAN-SF benchmark. Federated flare forecasting itself predates this work (Fu et al., 2023 preprint); what the documented literature search of Section 2.4 supports is that this is the first federated evaluation of SWAN-SF, the first instance-level provenance audit of the cleaned export (the leakage phenomenon itself is documented prior art — the benchmark’s creators warned about random-split overlap in 2019, and the cleaned release prescribes temporally-preceding folds), and the first leakage-free re-evaluation of federated flare prediction under the benchmark’s intended protocol. The study is a benchmark audit: six simulated clients (neutral labels, no real institution) hold disjoint Dirichlet-skewed shards of the public SWAN-SF benchmark, MLP clients train under FedAvg and FedProx with a clamped server-aware focal loss, and a frozen, leakage-audited evaluation protocol makes every selection decision on validation data. We make no deployment claim: the benchmark’s data is public and freely poolable, the downstream flare→CME→storm→GIC chain is far from the quantity modelled and is not evaluated, and the calibration analysis finds no validation-fit decision layer that survives the prevalence shift.

To keep the summary honest, we classify the paper’s claims into three levels, tiered against the leakage-free fold of Section 6.1. **Demonstrated (leakage-free, single-pass, frozen protocol):** the cleaned SWAN-SF export’s same-partition train/test pairing shares instances (100% of flaring test windows; 85.7–90% TimeGAN-synthetic training positives), and the benchmark’s intended temporally-preceding protocol removes the overlap; on that protocol all arms generalise (ROC-AUC 0.906–0.978); the in-partition federated-versus-centralised gap does not survive (FedProx 0.976/0.308 vs architecture-matched centralised MLP 0.973/0.382); the threshold-transfer and calibration failures persist (only XGBoost’s posteriors are usable); and event-level detection is essentially solved for every arm (64–65 of 65 events, 23.4 h median first-alert lead—a label-boundary quantity at this 24-hour label horizon, not a usable warning time, Table 5) while the false-alarm burden separates the arms by roughly 7×. Retraining the identical frozen protocol on the RAW, unbalanced benchmark (natural 2.05% training prevalence, no rebalancing, no synthetic positives, the release’s imputation and normalisation reproduced in-pipeline with train-only parameters and verified against the released export on the audit-matched windows) then separates the arms decisively: the centralised models are substrate-robust (logistic regression unchanged at 0.978; centralised MLP PR-AUC rises to 0.445), XGBoost’s minority-class precision was inflated by the synthetic positives (PR-AUC 0.457→0.329), and every federated MLP arm degrades sharply (FedProx 0.930/0.149; FedAvg detects zero of 65 events — 38 of 65 (58.5%) at 0.28 false-alarm windows per day under the diagnostic’s batch-statistics treatment, the only one with a usable validation operating point, Section 6.3; SCAFFOLD, executed here for the first time on a real substrate, fails at 0.768/0.057, while the same implementation’s LSTM instance recovers to 0.970/0.294 — both cells measuring the sign-inverted, cold-start variant disclosed in Section 4, not the reference algorithm). The architecture qualifies the ending: the identical frozen protocol with the SolarLSTM, executed on the owner GPU, restores the federated arms substantially (FedAvg-LSTM 0.958/0.305 where the FedAvg MLP fell to

0.875/0.056; FedProx-LSTM 0.970/0.404, above its own pooled comparator 0.962/0.314 and XGBoost’s raw PR-AUC, though still below logistic regression 0.978/0.448, and at event level the substrate’s highest detection rate, SCAFFOLD-LSTM’s 55/65 at 0.4 false-alarm windows per day, bought at a higher alarm budget)—and the architecture-only control completes the re-attribution: removing the MLP’s three BatchNorm layers and changing nothing else restores FedAvg to 0.963/0.366 and FedProx to 0.975/0.345, at or above the sequence arms’ ranking and with FedProx above the centralised MLP on ROC-AUC (Section 6.4)—so the balanced-substrate standing is a property of BatchNorm under federation, not of federation per se and not of the flattened-feature representation as such, and no federated arm exceeds the best centralised baseline on any substrate’s window-level ranking (the SCAFFOLD-LSTM event-level detection rate is the one priced exception, achieved at roughly twice the centralised arms’ false-alarm burden, and itself an implementation-case measurement of the five departures disclosed in Section 4). **Demonstrated in-partition (protocol diagnostics, disclosed overlap):** the shipped evaluation path separates the FL algorithms—FedProx with $\mu=0.01$ is above FedAvg on PR-AUC on five of five seeds (one-sided Wilcoxon $p = 0.031$) with $2.7\times$ narrower ROC-AUC variance, the FedProx global model attains 97.7% of the centralised XGBoost reference’s in-partition ROC-AUC (0.954 vs 0.977), and the shipped-eval FedAvg drift does not occur at $\alpha=5.0$ —but the BatchNorm-transport diagnostic and the no-BatchNorm control re-attribute this separation: the identical FedAvg weights hold 0.951–0.953 test ROC-AUC under correct normalisation at every round, the saturated all-positive validation monitor ($F_1 = 2p/(1+p)$) restores a collapse-round checkpoint, and without BatchNorm layers neither algorithm diverges and the gap shrinks to at most 0.036 ROC-AUC (Sections 6.11–6.12). Also demonstrated in-partition: per-client SMOTE is inactive on this benchmark variant (zero synthetic samples); and both the centralised and federated models concentrate attribution on current helicity, flux-emergence proxies, and Lorentz-force terms. **Supported with caveats:** one FedAvg cell with distribution-aware aggregation exceeds the FedProx headline ROC-AUC, recall-at-FPR figures are window-level statistics over correlated windows, five-seed confidence intervals quantify seed variability only, and the leakage-free fold is a single fold and seed (fold-replication remains queued owner-side compute; the raw-substrate seed replication is discharged in Section 6.6). Two formerly listed caveats were resolved by executed re-runs: the optimisation-budget confound (the budget-matched 500-epoch re-run of Section 5.5 shows the in-partition FedProx gap survives parity, though the gap itself is an in-partition artifact) and the within-distribution client holdouts (the untouched-holdout evaluation of Section 6.16 shows the federation benefit is larger under clean estimates). **Not demonstrated:** cryptographic privacy (secure aggregation and differential privacy are implemented but their accuracy cost is unmeasured; poisoning defence is absent), leakage-free multi-seed/ablation replication and the remaining temporally-preceding folds (one fold executed here), deployment-grade calibration (under the prior shift the neural posteriors are worse than climatology on Brier and ECE, and the executed five-arm comparison shows no validation-fit calibrator rescues them), neural-arm false-alarm reduction at event level, the downstream GIC/SCADA pipeline, the cleaned-fold pairing of the sequence arms, and fold-replication breadth; the SCAFFOLD-LSTM arm, the seed-43 replication, and the natural-prevalence balancing ablation queued in earlier revisions have since been executed (Section 6.6: the four-arm federation matrix is complete, every arm’s ROC-AUC replicates within ± 0.005 , and per-client SMOTE is a clean negative). An independent re-execution from the public dataset artefacts — byte-verified against the frozen SHA-256 manifest by a committed, executable verification script (20/20 files; Section 5.6) — reproduces every in-partition headline metric exactly; eleven items of the re-run programme—the frozen-FPR report, the calibration arms, the untouched holdouts, the budget-matched and $\alpha=5.0$ re-runs, the raw-metadata provenance audit, the event-level/leakage-free evaluation on the temporally-preceding fold, the raw-substrate re-training with the SCAFFOLD arm, the federated LSTM arms, their seed-43 replication, and the natural-prevalence SMOTE ablation—have since been executed with results in Section 6;

the remaining items are listed in Section 8, each with the implemented-and-tested protocol that would upgrade the claim.

The performance message is correspondingly conditional. Matching the same-architecture pooled model’s ranking is, on the leakage-free protocol, a property the pooled reference shares rather than a federation advantage, and on the raw unbalanced benchmark that parity is BatchNorm-conditional: every federated configuration carrying BatchNorm is strictly dominated by simple centralised baselines, while every BatchNorm-free configuration reaches the centralised MLP’s ranking — the best sequence arm (FedProx-LSTM, 0.970/0.404) comes within 0.008 ROC-AUC and 0.044 PR-AUC of the best centralised baseline, and the no-BatchNorm FedProx MLP (0.975/0.345, Section 6.4) within 0.003/0.103, above the centralised MLP’s ROC-AUC. No federated arm exceeds the best centralised baseline on any substrate, so the case for federating this problem rests on methodology (controlled heterogeneity studies, parameter-only transport), not on predictive performance. For the machine-learning community the study is a caution and a demonstration on two levels. Instance-level provenance audits can invert the conclusions drawn from a widely used preprocessed benchmark. And evaluation stacks have failure modes of their own: parameters-only weight transport, thresholded monitors, and strict-inequality checkpoint selection are each mundane choices whose composition manufactured a seed-reproducible algorithm gap that survived three revisions of this paper—the cross-silo regime of the physical sciences has failure modes—prior mismatch, layered imbalance handling, optimiser–variance-reduction interactions—that the medical and mobile FL literature does not fully anticipate, and that are worth studying in their own right.

A Configuration Evolution of the Released Framework

The released framework reached its present form through six documented version iterations (v2.1–v2.6), each responding to a diagnosed failure mode. We record this evolution because reproducibility of *process*—not only of final numbers—is what future deployments in other domains will actually need, and because several of the fixes correct subtle interactions that a clean-slate reimplementaion would likely rediscover the hard way.

Table 17: Documented stability fixes and protocol revisions across framework versions (v2.1–v4.11), as recorded in the released source.

Version	Diagnosed failure mode	Fix
v2.1	LSTM models require 3D input while baselines require 2D; a single data path served neither correctly.	Dual data path: 2D statistical features for MLP and baselines; raw 60×24 tensors for LSTM.
v2.2	F_β -optimal thresholds were computed but not applied; tables and confusion matrices reflected the default threshold.	Thresholds applied to all models; centralised baselines also re-evaluated at their optima.
v2.3	SCAFFOLD NaN every round: <code>std()</code> with Bessel correction on size-1 tensors divides by zero, poisoning control variates; DA-FL cap of $1/N$ starved large clients; focal α drifted to 0.95.	Population std for variate clipping; DA-FL cap $2/N$; alpha clamp tightened.
v2.4	Dynamic alpha scaling (up to $3\times$) re-inflated positive-class weight, undoing the focal- α fix; clients over-predicted flares (recall ≈ 1.0 , precision ≈ 0.03).	Square-root-damped imbalance factor capped at $1.5\times$; progressive factor narrowed to 0.9–1.1.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v2.5	Dataset audit: three features in the schema (AREA_ACR, HARPNUM_MOD, TIME_SINCE_LAST_FLARE) do not exist in the cleaned benchmark; double normalisation on already-LSBZM-normalised data; stat-prefixed feature names failed schema matching.	Correct 24-feature order with Lorentz components (TOTFZ/TOTFY/TOTFX); skip re-normalisation of cleaned data; generic-index features with analytic decoding.
v2.6	$\alpha=0.5$ partitions produced 3%–99% local flare rates as then observed (0.3%–99.8% re-measured under the frozen seed in v3.9), destabilising aggregation; 3 local epochs starved LSTM temporal learning; mixup on 3D sequences created unrealistic interpolations.	Dirichlet $\alpha \rightarrow 1.0$; local epochs $3 \rightarrow 10$; learning rate $10^{-3} \rightarrow 5 \times 10^{-4}$; mixup and SCAFFOLD disabled pending rework.
v3.0	Evaluation-protocol audit: thresholds selected on the test set; stale accuracy reported for federated models; non-disjoint Dirichlet shards; SMOTE dead code; DA-FL applied to both algorithm arms; one configuration constant scattered across modules.	Frozen protocol: validation-only selection, single-pass test, runtime leakage audit, disjoint-by-construction partitioning; explicit ablation switches (aggregation, SMOTE, loss, calibration); single <code>config.py</code> source of truth; centralised-MLP architecture control added.
v3.0.1	Single-seed headline lacked statistical support; results not machine-readable end-to-end.	Frozen-protocol rerun on the full cleaned benchmark (seed 42, 50 rounds); five-seed study with Wilcoxon tests; 24-cell ablation grid; μ/α sweep; SMOTE-validity, interpretability, and communication artefacts; SHA-256 data manifest; 74-check integrity suite.
v3.1	External review: region overlap across train/validation; balanced-prevalence validation cannot support threshold transfer or calibration; client-level evaluation optimistic; centralized-FL optimisation budgets unstated; window-level metrics read as event metrics; client labels implied institutional participation.	Protocol extensions, all unit-tested and queued for the full-data re-run: region-disjoint splits, natural-prevalence validation with validation-frozen FPR operating points, untouched client holdouts, training-budget audit, event-level metrics, calibration-comparison harness; neutral client labels; threat-model table; documented literature-search protocol.
v3.2	Second external review (25 findings): scope/title mismatch, premise overreach, stale v2.x text, claim taxonomy, budget confound, calibration honesty, communication-cadence ambiguity.	Data-locality scope title; claims taxonomy (verified / benchmark-only / not demonstrated); threat-model and budget tables; ECE column; literature-search protocol; lit-search artefacts; three regenerated neutral figures; response map docs/REVIEW2_RESPONSE.md.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v3.3	Reproducibility unproven end-to-end; calibration arms, frozen-FPR realisation, and untouched holdouts still unexecuted on real data.	Independent re-execution from the public Cleaned SWAN-SF release (SHA-256-verified against the data manifest): every headline test metric of all six models reproduced exactly (bit-stable CPU protocol); calibration-arm comparison executed on the trained models (negative result: no validation-fit arm survives the $26\times$ prior shift); untouched-holdout client evaluation executed (local 0.983 / FedAvg 0.866 / FedProx 0.988 mean PR-AUC); frozen-FPR realised test FPR reported; budget-matched 500-epoch centralised re-run executed (the shipped reference was undertrained; at parity 0.917/0.219 vs FedProx 0.954/0.307—gap survives); $\alpha=5.0$ promotion executed (negative; $\alpha=1.0$ retained); fig:clients regenerated from the clean protocol; runners made cache-reusing and round-resumable.
v3.4	Same-partition protocol: train/test instance overlap unknown and assumed absent; event-level and leakage-free evaluation impossible without raw metadata; four re-run items queued on that dependency.	Raw SWAN-SF benchmark obtained (Harvard Dataverse) and every cleaned window aligned to its raw instance via normalization-invariant matching (98.7–100% verification, 100% label agreement): instance-sharing quantified (test exports contain all raw instances; 100% of flaring test windows are training windows; 85.7–90% of training positives TimeGAN-synthetic); leakage-free fold executed under the benchmark’s intended temporally-preceding protocol (train P1–4 $\rightarrow$ test P5, round-resumable runner <code>run_partition_disjoint.py</code>): central-vs-FL gap disappears (FedProx 0.976/0.308 vs centralised MLP 0.973/0.382), FedProx ROC stabilisation over FedAvg survives (0.976 vs 0.906, checkpoint-independent), PR ordering reverses (FedAvg 0.370 above FedProx 0.308); calibration/threshold-transfer failures persist; event-level evaluation executed on 65 physical M/X-class events (detection 98.5–100%, XGBoost false-alarm burden $\sim 7\times$ lower than neural arms, 23.4 h median lead to flare peak); claims re-tiered, in-partition results disclosed as such.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v3.5	Raw-substrate retraining (re-run item 9) and the SCAFFOLD arm (item 8, half) still queued; the release's imputation/normalisation stages not reproducible in-pipeline.	Both executed (<code>experiments/raw_substrate.py</code> , <code>run_raw_substrate.py</code> , <code>run_event_level_raw.py</code>): FPCKNN-style imputation and LSBZM-style normalisation reproduced from the release paper's description with train-only parameters and verified against the released export on the audit-matched windows (median Spearman $\rho = 0.898$ at observed-nonzero positions; release is an exact monotone image of raw, $\rho = 1.000$; R-value zero mass re-imputed by the release but preserved here); frozen protocol retrained on raw P1–4 at natural 2.05% prevalence, single-pass test on raw P5: centralised arms substrate-robust (LR 0.978/0.448, XGB 0.974/0.329, MLP 0.971/0.445), federated arms collapse (FedAvg 0.875/0.056, 0/65 events; FedProx 0.930/0.149; SCAFFOLD 0.768/0.057); event-level: central arms 61.5–80% detection at 0.2–0.3 FA windows/day; claims re-tiered across three substrates.
v3.6	Federated LSTM arm (item 8, second half) unexecuted on owner-side compute; raw-substrate collapse established only for the MLP arms, leaving its architecture dependence open.	LSTM arms executed on the owner GPU (RTX 4060, 2.6 h; <code>run_federated_lstm.py</code> , <code>outputs/raw_lstm_eval.json</code> ; substrate identity asserted): centrally the LSTM is the weaker architecture (0.962/0.314 vs 0.971/0.445) but the federated sequence arms do not collapse—FedAvg-LSTM 0.958/0.305, FedProx-LSTM 0.970/0.404, above its pooled comparator and XGBoost's raw PR-AUC, below logistic regression; the federated sequence arms' Brier beats the climatology floor; stabilisation observed on every executed substrate–encoder combination; CPU-only event-level runner released; claims re-tiered (the collapse is a flattened-feature-MLP property).
v3.7	Event-level behaviour of the LSTM arms unevaluated: the owner-GPU run predated the audit-metadata extraction, so the v3.6 release shipped the CPU runner staged but unexecuted against real probabilities.	Executed via <code>run_event_level_lstm.py</code> (torch-free, CPU, Windows <code>_aux</code> fallback added) over the GPU run's stored test probabilities and the audit metadata (65/65 GOES peaks matched, labels asserted equal): FedAvg-LSTM 35/65 events (53.8%, 0.2 FA windows/day) against its MLP counterpart's 0/65; FedProx-LSTM 49/65 (75.4%) at 0.4 FA windows/day with a 23.2 h median lead—FedProx-MLP's detection at $3.5\times$ lower false-alarm burden; no federated sequence arm reaches the 78.5–80% centralised detection frontier. Artefact <code>outputs/event_level_raw_lstm_p5.json</code> .

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v3.8	The remaining owner-side programme unexecuted: SCAFFOLD-LSTM absent (hardest MLP failure left open), sequence arms single-seed single-device, and the natural-prevalence SMOTE ablation unstudied.	Executed as one resumable owner-GPU queue (8.7 h, RTX 4060; queue driver <code>run_gpu_queue.py</code> ; artefacts validated on receipt by a 64-check audit: protocol fingerprints, substrate identity, embedded frozen-counterpart bit-matches, event-key reconciliation, log cross-checks): SCAFFOLD-LSTM 0.970/0.294 window-level (vs MLP counterpart 0.768/0.057) and 55/65 events (84.6%) at 0.41 FA windows/day—the substrate’s highest detection rate, above the centralised frontier at a priced alarm budget (both cells measuring the disclosed five-departure implementation, not the reference algorithm); seed-43 replication of all four arms (ROC within ± 0.005 on every arm, qualitative structure unchanged); per-client SMOTE ablation a clean negative (FedAvg 0.961/0.307 unchanged, FedProx 0.946/0.288, -0.116 PR-AUC). Artefacts <code>raw_lstm_scaffold.json</code> , <code>raw_lstm_seed43.json</code> , <code>raw_lstm_smote.json</code> and their event-level JSONs; Tables 10–11 extended.
v3.9	External review identified the “FedAvg diverges” narrative as suspicious; the shipped checkpoints’ BatchNorm layers carried <code>num_batches_tracked=0</code> —running statistics never transported. Diagnostic: same weights, four BN treatments (<code>bn_diagnostic.json</code>); controlled no-BN re-run (<code>nobn_control.json</code>); saturation of the validation monitor derived in closed form; figure 2 regenerated from the realised partition (was an independent Dirichlet draw); reference [23] corrected; Fu et al. [11] cited, retiring first-to-federate claims; abstract compressed 781→257 words.	Interpretation frozen (<code>docs/INTERPRETATION_FREEZE.md</code>) before the rewrite; the divergence re-attributed to the evaluation path; Sections 6.11–6.12 added.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.0	Third external review (dual-chaired forensic dossier): the numbers are faithful but the provenance <i>meta-claims</i> are not substantiated — no verification code behind “byte-verified”, a hand-typed reference dict injecting a stale number into tab:leakfree, an internally inconsistent audit-JSON totals block (mixed row/unique-instance bases), an empty machine-written GPU ledger, four irreconcilable test counts, a wrong LSTM communication total, eleven metadata-defective bibliography entries, uncited prior art for both the leakage warnings and the BatchNorm mechanism, and a grid/GIC framing the paper’s own boundary document contradicts.	The evidence edition: <code>verify_manifest.py</code> (20/20 SHA-256, hard pipeline gate); reference dicts read from artefacts at run time and both JSONs regenerated; audit totals recomputed on dual bases with a committed recomputation test; queue logs committed and the ledger filled with artefact hashes; one battery runner, 214/214 checks; the Leka et al. verification apparatus (TSS/HSS/BSS/reliability/persistence/climatology) added to every operating point; bibliography repaired field-by-field against arXiv/DBLP/ADS; FedBN-lineage and prior-warning citations added; grid/GIC framing deleted from title, abstract, and operational claims; title narrowed to “benchmark audit”.
v4.1	Re-review of v4.0 (Dossier R-FS9-R1, verdict upgraded to major-revision-narrow-scope): the evidence edition’s own flagship additions arrived broken — a persistence baseline whose pool-row “time order” is temporally scrambled within active regions (stored TSS 0.348 matched no valid definition; the true label-inertia floor is 0.967 and no arm beats it), three aggregation bugs nulling the LSTM and frozen-FPR coverage, four NaN literals in the metrics JSON, a by-hand Poisson interval $\sim 2\times$ too narrow, five of six newly added bibliography entries carrying author lists from memory under a re-verification header, a broken cross-reference on page one, a figure whose labels contradicted its neutrality caption, and stale custodian/grid passages in the problem statement.	The corrected apparatus: persistence baselines recomputed on true time order with BOTH floors reported (label inertia 0.967; 24-hour-lag 0.418) and Section 6.1 rewritten around them (at TSS, XGBoost 0.848 and logistic regression 0.489 are the only arms above the lagged floor; at HSS none is); the two aggregation paths fixed (ten LSTM arms and 24 frozen-FPR cells now carry metrics) and the JSON regenerated as strict JSON, content-exact across the artefact-derived blocks; exact chi-square (Garwood) Poisson intervals; the six bib entries replaced with records read from arXiv/DBLP/IEEE/IOP; the <code>sec:limits</code> reference fixed; Figure 7 and the partition figure regenerated with plain A–F labels; the problem statement’s custodian and grid-protection passages replaced with the honest public-data simulation framing; the manifest environment claim matched to artefacts; a leakage-gate failure-path test added (battery 219/219).

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.2	Re-review of v4.1 (Dossier R-FS9-R2, verdict: minor revision, conditional acceptance): all previous blocking items verified fixed, both persistence floors independently re-implemented by the panel to the last digit, and the rebuilt bibliography re-checked against live Crossref/arXiv records; one new substantive error — the v4.1 correction itself had introduced a false exclusivity, “XGBoost alone clears the 24-hour-lagged floor,” refuted by the paper’s own artefact (logistic regression clears it at TSS 0.489; at HSS no arm clears, XGBoost $0.201 < 0.434$) — plus five minors: an event-level false-alarm range whose lower endpoint (8.6) belonged to logistic regression rather than the neural arms, two execution environments conflated in one parenthetical (the Python 3.12 protocol run vs. the battery header’s 3.13.5), a “bit-exact” overstatement (the metrics artefact embeds wall-clock <code>elapsed_s</code> ; the artefact-derived blocks are content-exact), a response-letter arm miscount, and a stale docstring cadence.	The exclusivity clause corrected in all four places (abstract, Section 6.1, the v4.1 row above, response letter) to the artefact-backed statement — at TSS two arms clear the lagged floor, at HSS none does — and logistic regression’s TSS/HSS added to Section 6.1’s first finding; the false-alarm range re-attributed (11.6–12.7 for the neural arms); the environment parenthetical split; “bit-exact” softened to content-exact with block B’s phase-cache dependency stated; the lag-definition sensitivity sweep committed (<code>experiments/run_lag_definition_sweep.py</code> → <code>outputs/lag_definition_sweep.json</code> : seven point-lag variants land at TSS 0.392–0.418, two arms clearing each; the lookback-flare variant scores 0.964, label inertia under another name); the docstring cadence corrected to 60 minutes; the battery re-verified in the current (torch-less) verification environment at 195 passed / 0 failed with the 24 torch-dependent checks skipped — the same delta the panel’s own fresh run showed — and the committed v4.1 log (219/219) retained as the full-battery record (<code>logs/test_battery_v4.2.log</code>).
v4.3	Re-review of v4.2 (Dossier R-FS9-R3, verdict: accept, conditional on four clerical items): all seven third-round findings verified resolved (one out-of-band), zero new substantive errors, and the lag-definition sweep re-executed byte-for-byte in the referee’s torch-less environment; a four-item clerical residue register — a response-letter erratum claim unverifiable from the repository, a mistitled battery log, the new sweep artefact lacking battery coverage, and a missing ledger hash — stood between the revision and closure.	The corrected response letters committed under <code>docs/response_letters/</code> (the R-FS9-R1 errata, the R-FS9-R2 and R-FS9-R3 letters, with a README), so every letter-level ledger claim is verifiable in-repo; the battery log title made version-bearing (a <code>BATTERY_VERSION</code> constant; the v4.2 log retitled in place) so the title cannot trail the release again; <code>tests/test_lag_sweep_artifact.py</code> added — 29 checks validating the sweep artefact against <code>standard_metrics.json</code> at run time (strict-JSON structure, the shipped rule vs block C at 10^{-9} , the clears-the-floor lists recomputed from block A for every definition, point-lag invariance, and the lookback variant’s disclosed n), so a silent edit that broke cross-artefact agreement now fails the battery; the sweep digest appended to the RUNLOG ledger; an environments manifest added at <code>docs/ENVIRONMENTS.md</code> (the panel’s recorded observation); the battery re-verified at 224 passed / 0 failed (the v4.2 count of 195 plus the 29 new checks; the same 24 torch-dependent skips; <code>logs/test_battery_v4.3.log</code> ; the canonical v4.1 log retained untouched).

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.5	Post-acceptance full-body re-audit (Dossier R-FS9-R5, verdict: accept sustained, composite rating 7.9/10): every bibliography entry verified field-by-field against publisher records, eleven exact-phrase plagiarism probes clean, both novelty firsts re-derived independently, the headline numbers reality-checked against the operational forecasting literature, and the never-audited code paths read line-by-line; seven new non-blocking findings registered (R7-1–R7-4, C1–C3), discharged as one-line errata. The intervening v4.4 was a repository-only errata batch (Dossier R-FS9-R4) that did not touch the manuscript.	The bibliography’s three defects fixed (a one-letter author misspelling restored to the Crossref record, a twelve-of-twenty-eight consortium roster marked “and others”, and two orphaned BibTeX lines deleted); the SCAFFOLD protocol disclosure completed (optimiser SGD with momentum 0.9 at half the global learning rate, and the arm’s pre-v4.5 focal-prevalence asymmetry stated plainly, with the call site aligned so all three federated arms share one loss definition); the budget artefact’s federated batch size regenerated to the loader-true value it had always trained at; the battery’s default-path log un-shipped and gitignored so a fresh clone’s first documented run succeeds; dead configuration constants deleted; the repository front door refreshed to the v4.5 state.
v4.6	Submission-manuscript audit (Dossier R-FS9-R6, verdict: major revision on the submission package, 5.9/10; the repository side verified with all eleven v4.3–v4.5 discharge items confirmed and the battery reproduced at 224/0): the round audited, for the first time, the hand-prepared journal manuscript beside the repository and found it regressed to the v3.9-era state (eleven defective bibliography entries, the floor apparatus and both SCAFFOLD disclosures absent, six prior-art citations dropped, two untraceable constants) — the fourth instance of the hand-prepared-artefact drift species; it also found a critical import regression at HEAD (a still-imported constant deleted at v4.5), an undisclosed sign inversion in the SCAFFOLD control-variate update, a false warm-start disclosure, and two comparator contradictions.	The import graph repaired and guarded by a new AST-based battery test that resolves every local import in every environment; the SCAFFOLD record corrected to a five-departure enumeration (the sign inversion disclosed as the fourth departure, the cold start every published operating point actually used as the fifth, the warm-start claim withdrawn) with the arm’s comparisons restated as implementation-case evidence; the budget table’s scheduler cell corrected (centralised arms train at constant rate); the focal-prevalence plumbing extended to the centralised comparators with the asymmetry disclosed; the focal loss’s two remaining constants disclosed; the interpretability fallback made loud and recorded; block A’s denominators read from the artefacts with self-checks; the lit-search dates embedded; and the submission manuscript regenerated from the record (submission/ in the repository, generator + battery guard — the drift species ends here).

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.7	Regenerated-submission audit (Dossier R-FS9-R7, verdict: accept conditional, 7.4/10): all six v4.6 register items verified against the code and the compiled PDFs, the integrity battery reproduced in two independent fresh trees, the submission regeneration reproduced byte-for-byte, and the six new guards probed by fault-injection; one blocking clerical condition — the v4.5 bibliography errata had placed their explanatory comments <i>inside</i> two entries, where the bibtex parser silently swallowed every field after the author list, so references [7] and [22] rendered as bare author lists in both PDFs through two guarded compiles (the guards verified source byte-identity, which faithfully inherited the record’s own rendering defect).	The two comments moved outside the entries and both references render complete in the paper and the regenerated submission; a new rendered-bibliography check extracts the compiled PDFs and asserts the entry count, a year on every entry, and the title fragments and years of the pinned load-bearing entries; the six guards hardened per the panel’s recommendations (a manifest registry so new letters cannot opt out of the diff-manifest convention, crashed-guard detection in the battery runner, plain-import and module-deletion coverage in the import graph, the promised argmax-consistency check for the interpretability artefact, the pinned block-A raw-substrate prevalence, and asserted counts against descriptive drift).
v4.8	Post-closure content errata (Dossier R-FS9-R8 closed the cycle with verdict accept—sustained, 7.9/10; this revision folds in its non-blocking residual register C1–C6/T1–T2 and the owner-side content review’s E1–E3 conditions): the raw-substrate federated MLP arms were disclosed as evaluated through the same untransported-BatchNorm path the paper’s own diagnostic indicts in-partition — the MLP degradation is BN-evaluation-conditional until the checkpoint-level diagnostic runs (E1; the runner and its battery guard ship with this revision, execution queued owner-side; the LSTM arms carry no BatchNorm and are unaffected); the SCAFFOLD implementation-case qualifiers extended to every 55/65 mention (E2); the abstract reordered to lead its results claims with the corrected persistence finding (E3); one display erratum (the SCAFFOLD-MLP median lead is exactly 23.25 h, shown at full precision).	The register folded in: a non-zero exit is now a battery failure whatever its verdict lines (the PASS-then-die and SKIP-then-die escapes closed, tracebacks no longer suppressed, C1/T1), the letter registry is an explicit pinned map rather than a filename-version regex (C2), all 31 bibliography entries pinned by anchor + title + year with the pin count asserted (C3), the erratum line recorded for the v4.7 letter’s phantom “ade” restoration — a display-pipeline artefact, no restoration happened (C4), the central-criterion scan extended to provenance/ and data_manifest/ (C5), and the import graph enforces a pinned external whitelist so an unresolved import is an error, not an implicit external (C6).

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.9	Results-integration errata (RUNLOG ask #4 executed owner-side — the E1 raw-substrate BN diagnostic): the artefact of record <code>raw_bn_diagnostic.json</code> committed and sha-pinned, the diagnostic's outcome integrated (new Section 6.3 + Table 8), and the v4.8 BN-evaluation-conditional qualifier resolved <i>against</i> the BN explanation — the raw-substrate MLP deficit survives every treatment (pooled recalibration collapses all arms to 0.500–0.501; batch statistics leave the deficit standing), the opposite of the in-partition pattern; the abstract clause and the limitations qualifier updated; the same-weights counterfactual disclosed as unavailable (the v3.x-era checkpoints were never persisted — the diagnostic ran over a fresh seed-42 replication whose gate honestly recorded MISMATCH, frozen and battery-pinned).	Environment hardening from the owner's v4.8 battery run (316/6): the frozen artefact trees are now <code>-text</code> in <code>.gitattributes</code> so every platform checks them out byte-identical (a Windows autocrlf checkout had failed the interpretability sha pin on a content-identical record); the raw-BN guard re-pinned from the same-weights happy path to the disclosed MISMATCH state; <code>test_raw_lstm</code> exits through a disclosed win32 teardown guard (the 0xC0000409 interpreter-finalization fastfail the ask #3 queue children already exhibited — a battery module must exit 0 after a green run); <code>test_fl_smoke</code> 's CUDA-tensor <code>.numpy()</code> sites fixed.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.10	Register-integration errata (the Round-13 master change register’s paper-side phase; the repo-only v4.9.2–v4.9.4 errata chain shipped the instruments, the owner executed them, and this revision integrates the verdicts): A1 EXECUTED, verdict <code>arm_b_validated</code> — arm B applied to the centralised raw MLP preserves the ranking (-6.8×10^{-5} ROC-AUC, $+4.9 \times 10^{-3}$ PR-AUC against own buffers, both gates match; <code>arm_b_central_sanity.json</code> committed and battery-checked) — so Table 8’s B column is re-framed as the <i>validated proxy</i> for the pooled-statistics transport fix and the faithful arm-D re-run is not required for the reading (A2, the re-frame branch); A3 EXECUTED — the no-BatchNorm raw-substrate federation (<code>raw_nobn_eval.json</code>, new Section 6.4 + Table 9) re-attributes the raw-substrate federated MLP collapse to BatchNorm-under-federation (no-BN FedAvg 0.963/0.366, FedProx 0.975/0.345, the latter above the centralised MLP on ROC-AUC), and contribution 5 is rewritten accordingly — the encoder-conditional attribution is withdrawn; A4 the abstract’s closure clause scoped to the executed attribution, the “TSS 0.404” label corrected to PR-AUC, the internal version tag dropped; B1 the degenerate checkpoint rule’s round disclosed (the strict-inequality rule ships the final round on an all-zero history), the new arms’ val-ROC sensitivity rows and per-round test trajectories integrated, and Section 9’s “zero of 65” re-qualified with the arm-C 38/65 companion cell; B2 the implementation-case qualifier extended to Table 8’s SCAFFOLD row; B3 the abstract’s persistence sentence reordered (lagged persistence leads, label inertia second); B4 <code>raw_substrate_rerun.json</code> named in Section 6.3 and sha-pinned by the battery; B5 the abstract’s manifest sentence separated (result artefacts vs the 20-file data manifest).	Apparatus: the page-count pins bumped with the disclosed recompile (the B7 convention); <code>raw_substrate_rerun.json</code> sha-pinned in <code>test_raw_bn_diagnostic.py</code> (B4); the submission regenerated from the record (<code>tools/build_submission.py</code>, its T2 metric-name assertion re-pinned to the corrected PR-AUC label); the letters registry carries this revision’s letter; battery version v4.10.

(continued on next page)

(continued from previous page)

Version	Diagnosed failure mode	Fix
v4.11	Pre-submission errata (Dossier R-FS9-R11 — the register-integration verification round: verdict accept-sustained 8.4/10, the Round-13 condition discharged, nothing blocking; the panel’s punch list): the past-zero residual re-rounded to -0.004 (two sites, Section 6.4 and the gap list of Section 6.7); the 6.3 sentence’s “both” antecedent corrected to arm A alone; the bare “zero of 65” sites companioned with the arm-C 38/65 cell (intro contribution 5, the event-level table’s new arm-C row, Section 6.2); the abstract’s “should reconstruct” restored and both “ranking” claims metric-scoped to ROC-AUC (the no-BN PR ordering inverts between the arms); 6.5 re-titled from the encoder rescue to the BatchNorm-free rescue; the FedAvg trajectory’s final-step $+9 \times 10^{-6}$ uptick disclosed where “monotonically” stood; the 6.7 gap list’s sign convention unified (signed federated-minus-reference, every magnitude unchanged)	Apparatus (P1/P2/P7): the generator’s version stamp single-sourced from <code>BATTERY_VERSION</code> at generation time — the stale title-page string class is closed; raw-BN guard layer 5 pins both decisive artefacts by sha256 and the Table 9 region’s seven rows against live artefact-derived values (an in-range tamper of <code>raw_nobn_eval.json</code> or <code>arm_b_central_sanity.json</code> now fails the battery), and a source-comparison guard enforces the arm-B machinery’s byte-identity; submission regenerated (stamp v4.11); the letters registry carries this revision’s letter; battery version v4.11.

Three of these fixes deserve narrative emphasis because they encode non-obvious engineering knowledge. The SCAFFOLD failure (v2.3) is the most instructive: the NaN symptoms pointed nowhere near their cause, which was a statistics API detail—Bessel’s correction on single-element bias tensors—interacting with an algorithm whose correctness depends on exact control-variate arithmetic under fixed learning rates (AdamW’s adaptive steps further invalidate the update formula). The focal-alpha saga (v2.3–v2.4) illustrates layered imbalance handling drifting into degeneracy: SMOTE, class-weighted loss, and rate-aware aggregation all push the same effective prior, and unbounded adaptation at any one layer can silently dominate the others. The dataset audit (v2.5) is a reminder that in applied ML the schema *is* the experiment: three phantom features had been silently participating in the pipeline until an attribute-by-attribute audit against the benchmark’s actual storage order replaced them with the physically meaningful Lorentz-force components. The v3.x revisions add a fourth, blander lesson: protocol defects (test-set selection, optimistic client holdouts, unstated budgets) do not announce themselves in any loss curve—only an external audit found them—which is the argument for freezing the evaluation contract in code, as the released framework now does. The published artefacts, figures, and this paper’s numbers all post-date these fixes.

References

- [1] Rafal A. Angryk, Petrus C. Martens, Berkay Aydin, Dustin Kempton, Sushant S. Mahajan, Sunitha Basodi, Azim Ahmadzadeh, Xumin Cai, Soukaina Filali Boubrahimi, Shah Muhammad Hamdi, Michael A. Schuh, and Manolis K. Georgoulis. Multivariate time series dataset for space weather data analytics. *Scientific Data*, 7:227, 2020. doi: 10.1038/s41597-020-0548-x.
- [2] Azim Ahmadzadeh, Maxwell Hostetter, Berkay Aydin, Manolis K. Georgoulis, Dustin Kempton, Sushant Mahajan, and Rafal A. Angryk. Challenges with extreme class-imbalance

- and temporal coherence: A study on solar flare data. In *IEEE International Conference on Big Data (Big Data)*, pages 1423–1431, 2019. arXiv:1911.09061; v4.1: roster and published title verified against the IEEE / Google Scholar records (the v4.0 entry copied the 2020 dataset-paper roster).
- [3] Azim Ahmadzadeh, Berkay Aydin, Manolis K. Georgoulis, Dustin Kempton, Sushant Mahajan, and Rafal A. Angryk. How to train your flare prediction model: Revisiting robust sampling of rare events. *The Astrophysical Journal Supplement Series*, 254(2):23, 2021. doi: 10.3847/1538-4365/abec88. v4.1: roster, pages and DOI verified against the IOP / ADS records (the v4.0 author list was wrong).
 - [4] MohammadReza EskandariNasab, Shah Muhammad Hamdi, and Soukaina Filali Boubrahimi. Enhancing multivariate time series-based solar flare prediction with multifaceted preprocessing and contrastive learning. *arXiv preprint arXiv:2409.14016*, 2024.
 - [5] Monica G. Bobra and Sébastien Couvidat. Solar flare prediction using SDO/HMI vector magnetic field data with a machine-learning algorithm. *The Astrophysical Journal*, 798(2): 135, 2015. doi: 10.1088/0004-637X/798/2/135.
 - [6] K. D. Leka, Sung-Hong Park, Kanya Kusano, Jesse Andries, Graham Barnes, et al. A comparison of flare forecasting methods. II. benchmarks, metrics, and performance results for operational solar flare forecasting systems. *The Astrophysical Journal Supplement Series*, 243(2):36, 2019. doi: 10.3847/1538-4365/ab2e12. arXiv:1907.02905.
 - [7] Manolis K. Georgoulis, D. Shaun Bloomfield, Michele Piana, Anna Maria Massone, Marco Soldati, Peter T. Gallagher, Etienne Pariat, Nicole Vilmer, Eric Buchlin, Frederic Baudin, Andre Csillaghy, Hanna Sathiapal, et al. The flare likelihood and region eruption forecasting (FLARECAST) project: flare forecasting in the big data & machine learning era. *Journal of Space Weather and Space Climate*, 11:39, 2021. doi: 10.1051/swsc/2021023.
 - [8] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, pages 1273–1282, 2017. arXiv:1602.05629.
 - [9] Peter Kairouz, H. Brendan McMahan, Brendan Avent, Aurélien Bellet, Mehdi Bennis, Arjun N. Bhagoji, et al. Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1–2):1–210, 2021. arXiv:1912.04977.
 - [10] Nicola Rieke, Jonny Hancox, Wenqi Li, Fausto Milletari, Holger R. Roth, Shadi Albarqouni, et al. The future of digital health with federated learning. *npj Digital Medicine*, 3:119, 2020. doi: 10.1038/s41746-020-00323-1.
 - [11] Junfeng Fu, Jie Wan, Xin Huang, Ke Han, and Edison Peng. Federated transfer learning for solar flare forecasting [sic] — title as published. *Research Square preprint (not peer reviewed)*, 2023. doi: 10.21203/rs.3.rs-3134736/v1. Posted 10 July 2023; formatted for Solar Physics.
 - [12] Xiaoxiao Li, Meirui Jiang, Xiaofei Zhang, Michael Kamp, and Qi Dou. FedBN: Federated learning on non-IID features via local batch normalization. In *International Conference on Learning Representations (ICLR)*, 2021. arXiv:2102.07623.
 - [13] David H. Boteler. A 21st century view of the march 1989 magnetic storm. *Space Weather*, 17(10):1427–1441, 2019. doi: 10.1029/2019SW002278.

- [14] Dipankar Sarkar, Ankur Narang, and Sumit Rai. Fed-focal loss for imbalanced data classification in federated learning. *arXiv preprint arXiv:2011.06283*, 2020.
- [15] Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia Smith. Federated optimization in heterogeneous networks. In *Proceedings of Machine Learning and Systems (MLSys)*, volume 2, pages 429–450, 2020. arXiv:1812.06127.
- [16] Sai Praneeth Karimireddy, Satyen Kale, Mehryar Mohri, Sashank Reddi, Sebastian Stich, and Ananda Theertha Suresh. SCAFFOLD: Stochastic controlled averaging for federated learning. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, pages 5132–5143, 2020. arXiv:1910.06378.
- [17] Yanmeng Wang, Qingjiang Shi, and Tsung-Hui Chang. Why batch normalization damage federated learning on non-IID data? *IEEE Transactions on Neural Networks and Learning Systems*, 36(1):1692–1706, 2025. doi: 10.1109/TNNLS.2023.3323302. arXiv:2301.02982; v4.1: roster and TNNLS record verified against DBLP and IEEE Xplore (the v4.0 author list was wrong).
- [18] Rachid Guerraoui, Rafael Pinot, Geovani Rizk, John Stephan, and François Taiani. Overcoming the challenges of batch normalization in federated learning. *arXiv preprint arXiv:2405.14670*, 2024. v4.1: roster and title verified against the arXiv record; no “Heterogeneous” in the title and no additional authors (the v4.0 list was wrong).
- [19] Gonzalo Inaki Quintana, Laurence Vancamberg, Vincent Jugnon, Mathilde Mougeot, and Agnès Desolneux. BN-SCAFFOLD: Controlling the drift of batch normalization statistics in federated learning. *arXiv preprint arXiv:2410.03281*, 2024. v4.1: roster and title verified against arXiv / Semantic Scholar / Google Scholar records (the v4.0 author list was wrong).
- [20] Carolus J. Schrijver. A characteristic magnetic field pattern associated with all major solar flares and its use in flare forecasting. *The Astrophysical Journal Letters*, 655(2):L117–L120, 2007. doi: 10.1086/511857.
- [21] Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. Time-series generative adversarial networks. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 5508–5518, 2019.
- [22] Zeinab Hassani, Davud Mohammadpur, and Hossein Safari. Solar flare prediction using long short-term memory (LSTM) and decomposition-LSTM with sliding window pattern recognition. *The Astrophysical Journal Supplement Series*, 279(1):27, 2025. doi: 10.3847/1538-4365/addc73. preprint: arXiv:2507.05313.
- [23] Yue Zhao, Meng Li, Liangzhen Lai, Naveen Suda, Damon Civin, and Vikas Chandra. Federated learning with non-IID data. *arXiv preprint arXiv:1806.00582*, 2018.
- [24] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2999–3007, 2017. arXiv:1708.02002.
- [25] Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, and W. Philip Kegelmeyer. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16:321–357, 2002. doi: 10.1613/jair.953.
- [26] Hongyi Zhang, Moustapha Cissé, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations (ICLR)*, 2018. arXiv:1710.09412.

- [27] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. doi: 10.1162/neco.1997.9.8.1735.
- [28] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019. arXiv:1711.05101.
- [29] Ligeng Zhu, Zhijian Liu, and Song Han. Deep leakage from gradients. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 14774–14784, 2019. arXiv:1906.08935.
- [30] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pages 4765–4774, 2017.
- [31] Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.